

sunbase-assignment

September 27, 2023

0.1 Setup: Imports and Installations

```
[ ]: !pip install catboost
```

Collecting catboost

Downloading catboost-1.2.2-cp310-cp310-manylinux2014_x86_64.whl (98.7 MB)
98.7/98.7 MB

5.4 MB/s eta 0:00:00

Requirement already satisfied: graphviz in /usr/local/lib/python3.10/dist-packages (from catboost) (0.20.1)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from catboost) (3.7.1)

Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.10/dist-packages (from catboost) (1.23.5)

Requirement already satisfied: pandas>=0.24 in /usr/local/lib/python3.10/dist-packages (from catboost) (1.5.3)

Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from catboost) (1.11.2)

Requirement already satisfied: plotly in /usr/local/lib/python3.10/dist-packages (from catboost) (5.15.0)

Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from catboost) (1.16.0)

Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.24->catboost) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.24->catboost) (2023.3.post1)

Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (1.1.0)

Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (0.11.0)

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (4.42.1)

Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (1.4.5)

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (23.1)

Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (9.4.0)

Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (3.1.1)
Requirement already satisfied: tenacity>=6.2.0 in
/usr/local/lib/python3.10/dist-packages (from plotly->catboost) (8.2.3)
Installing collected packages: catboost
Successfully installed catboost-1.2.2

```
[ ]: import numpy as np
import pandas as pd
pd.set_option('display.max_columns', None)

## Data visualization
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib.ticker as ticker
import missingno as msno # missing values visualization

## Stats
from statsmodels.graphics.gofplots import qqplot
from scipy.stats import shapiro
from scipy.stats import norm
from scipy.stats import zscore

## Preprocessing
from sklearn.model_selection import train_test_split as tts
from sklearn.preprocessing import StandardScaler

## Classes Imbalance
from imblearn.over_sampling import SMOTE

## Models
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import ExtraTreesClassifier, RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from catboost import CatBoostClassifier
from sklearn.ensemble import VotingClassifier, StackingClassifier
from sklearn.ensemble import HistGradientBoostingClassifier

## Metrics
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix

## Time
import time
```

```

## Warnings
import warnings
warnings.filterwarnings('ignore')

# some more imports
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
from sklearn.model_selection import RandomizedSearchCV
from sklearn.impute import SimpleImputer

import os
for dirname, _, filenames in os.walk('/content/drive/MyDrive/internwork'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

```

```

[ ]: from google.colab import drive
drive.mount('/content/drive')

```

Mounted at /content/drive

1 1 Data Preprocessing:

Load the provided dataset and perform initial data exploration. Handle missing data and outliers. Prepare the data for machine learning by encoding categorical variables and splitting it into training and testing sets.

```

[ ]: df = pd.read_excel('/content/drive/MyDrive/internwork/
↳customer_churn_large_dataset.xlsx')
df

```

```

[ ]:

```

	CustomerID	Name	Age	Gender	Location \
0	1	Customer_1	63	Male	Los Angeles
1	2	Customer_2	62	Female	New York
2	3	Customer_3	24	Female	Los Angeles
3	4	Customer_4	36	Female	Miami
4	5	Customer_5	46	Female	Miami
...
99995	99996	Customer_99996	33	Male	Houston
99996	99997	Customer_99997	62	Female	New York
99997	99998	Customer_99998	64	Male	Chicago
99998	99999	Customer_99999	51	Female	New York
99999	100000	Customer_100000	27	Female	Los Angeles

	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	Churn
0	17	73.36	236	0
1	1	48.76	172	0

2	5	85.47	460	0
3	3	97.94	297	1
4	19	58.14	266	0
...
99995	23	55.13	226	1
99996	19	61.65	351	0
99997	17	96.11	251	1
99998	20	49.25	434	1
99999	19	76.57	173	1

[100000 rows x 9 columns]

```
[ ]: df.shape
```

```
[ ]: (100000, 9)
```

```
[ ]: # df.drop('CustomerID',axis=1,inplace=True)
# df.drop('Name',axis=1,inplace=True)
```

```
[ ]: df.tail(2)
```

```
[ ]:
      CustomerID      Name  Age  Gender  Location \
99998      99999  Customer_99999   51  Female    New York
99999     100000  Customer_100000   27  Female  Los Angeles

      Subscription_Length_Months  Monthly_Bill  Total_Usage_GB  Churn
99998                        20           49.25           434      1
99999                        19           76.57           173      1
```

Details of the dataset:

```
[ ]: print(df.head())
print(df.info())
print(df.describe())
```

	CustomerID	Name	Age	Gender	Location	\
0	1	Customer_1	63	Male	Los Angeles	
1	2	Customer_2	62	Female	New York	
2	3	Customer_3	24	Female	Los Angeles	
3	4	Customer_4	36	Female	Miami	
4	5	Customer_5	46	Female	Miami	

	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	Churn
0	17	73.36	236	0
1	1	48.76	172	0
2	5	85.47	460	0
3	3	97.94	297	1
4	19	58.14	266	0

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CustomerID                            100000 non-null  int64
1   Name                                  100000 non-null  object
2   Age                                   100000 non-null  int64
3   Gender                                100000 non-null  object
4   Location                              100000 non-null  object
5   Subscription_Length_Months            100000 non-null  int64
6   Monthly_Bill                          100000 non-null  float64
7   Total_Usage_GB                        100000 non-null  int64
8   Churn                                 100000 non-null  int64
dtypes: float64(1), int64(5), object(3)
memory usage: 6.9+ MB
None
```

	CustomerID	Age	Subscription_Length_Months \
count	100000.000000	100000.000000	100000.000000
mean	50000.500000	44.027020	12.490100
std	28867.657797	15.280283	6.926461
min	1.000000	18.000000	1.000000
25%	25000.750000	31.000000	6.000000
50%	50000.500000	44.000000	12.000000
75%	75000.250000	57.000000	19.000000
max	100000.000000	70.000000	24.000000

	Monthly_Bill	Total_Usage_GB	Churn
count	100000.000000	100000.000000	100000.000000
mean	65.053197	274.393650	0.497790
std	20.230696	130.463063	0.499998
min	30.000000	50.000000	0.000000
25%	47.540000	161.000000	0.000000
50%	65.010000	274.000000	0.000000
75%	82.640000	387.000000	1.000000
max	100.000000	500.000000	1.000000

Violin Plots to detect any possible outliers, but on plotting them we are able to see that there are negligible outliers

```
[ ]: import seaborn as sns
import matplotlib.pyplot as plt

# Assuming you have a DataFrame named 'df' with numerical columns
numerical_columns = df.select_dtypes(include=['number']).columns

# Set the number of columns per row in the subplot grid
num_plots_per_row = 2
```

```

# Calculate the number of rows needed based on the number of numerical columns
num_rows = len(numerical_columns) // num_plots_per_row
if len(numerical_columns) % num_plots_per_row != 0:
    num_rows += 1

# Create subplots
fig, axes = plt.subplots(num_rows, num_plots_per_row, figsize=(12, 5 *
    num_rows))

# Flatten the axes if there's only one row
if num_rows == 1:
    axes = axes.reshape(1, -1)

# Plot barplots for each numerical column
for i, column in enumerate(numerical_columns):
    row_idx = i // num_plots_per_row
    col_idx = i % num_plots_per_row

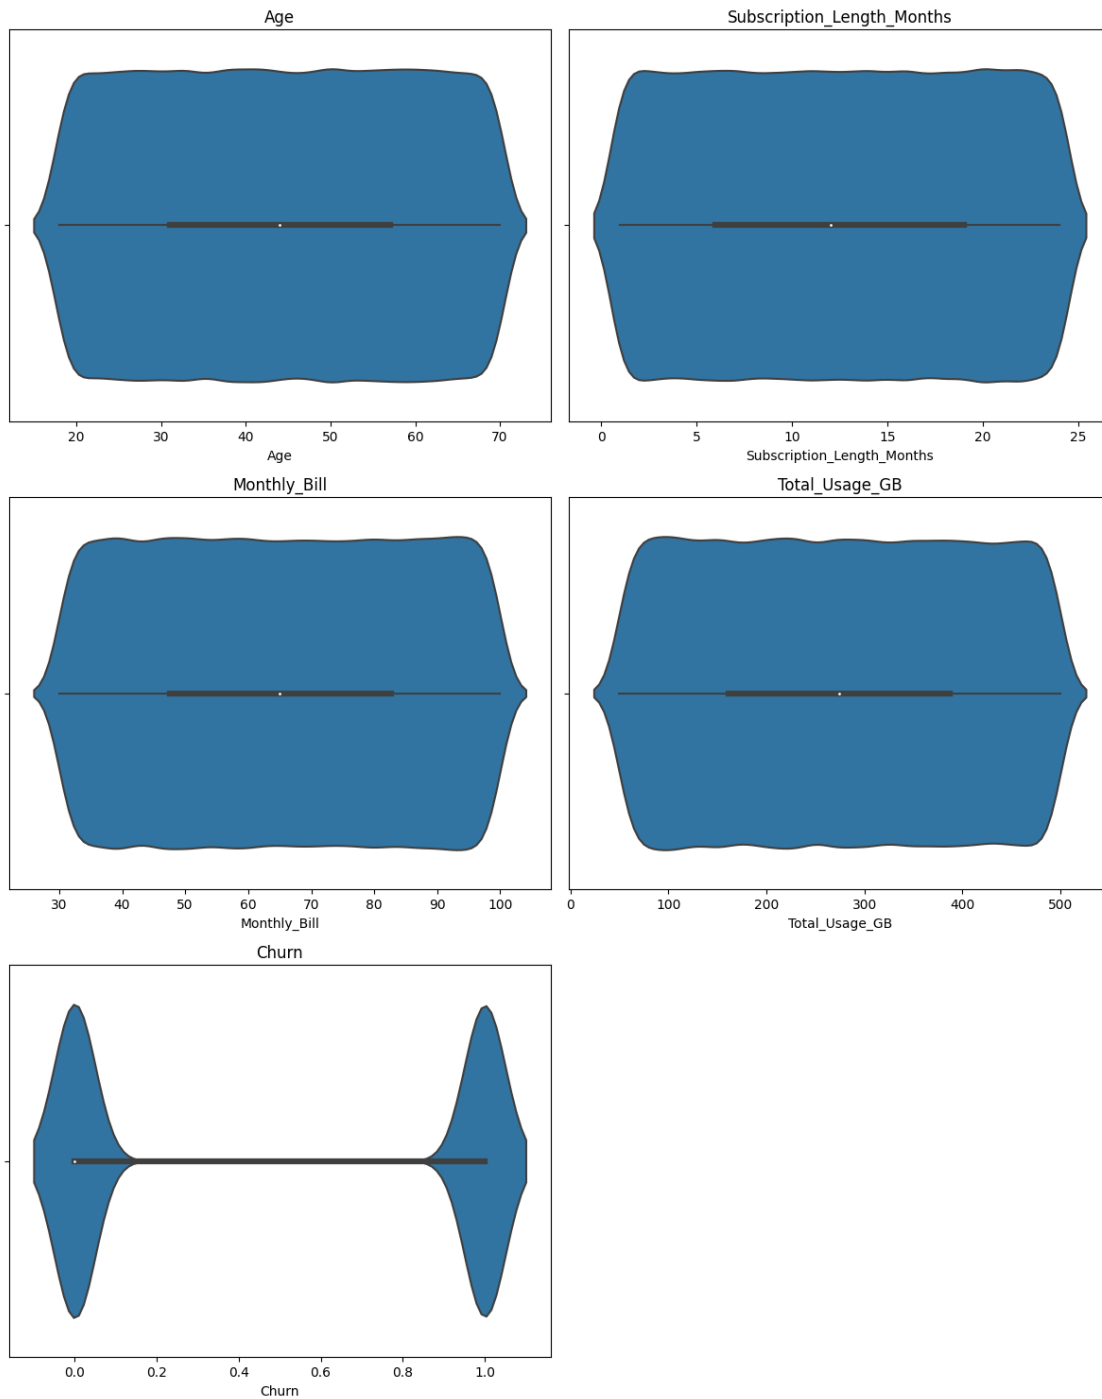
    ax = axes[row_idx, col_idx]

    sns.violinplot(data=df, x=column, ax=ax)
    ax.set_title(column)

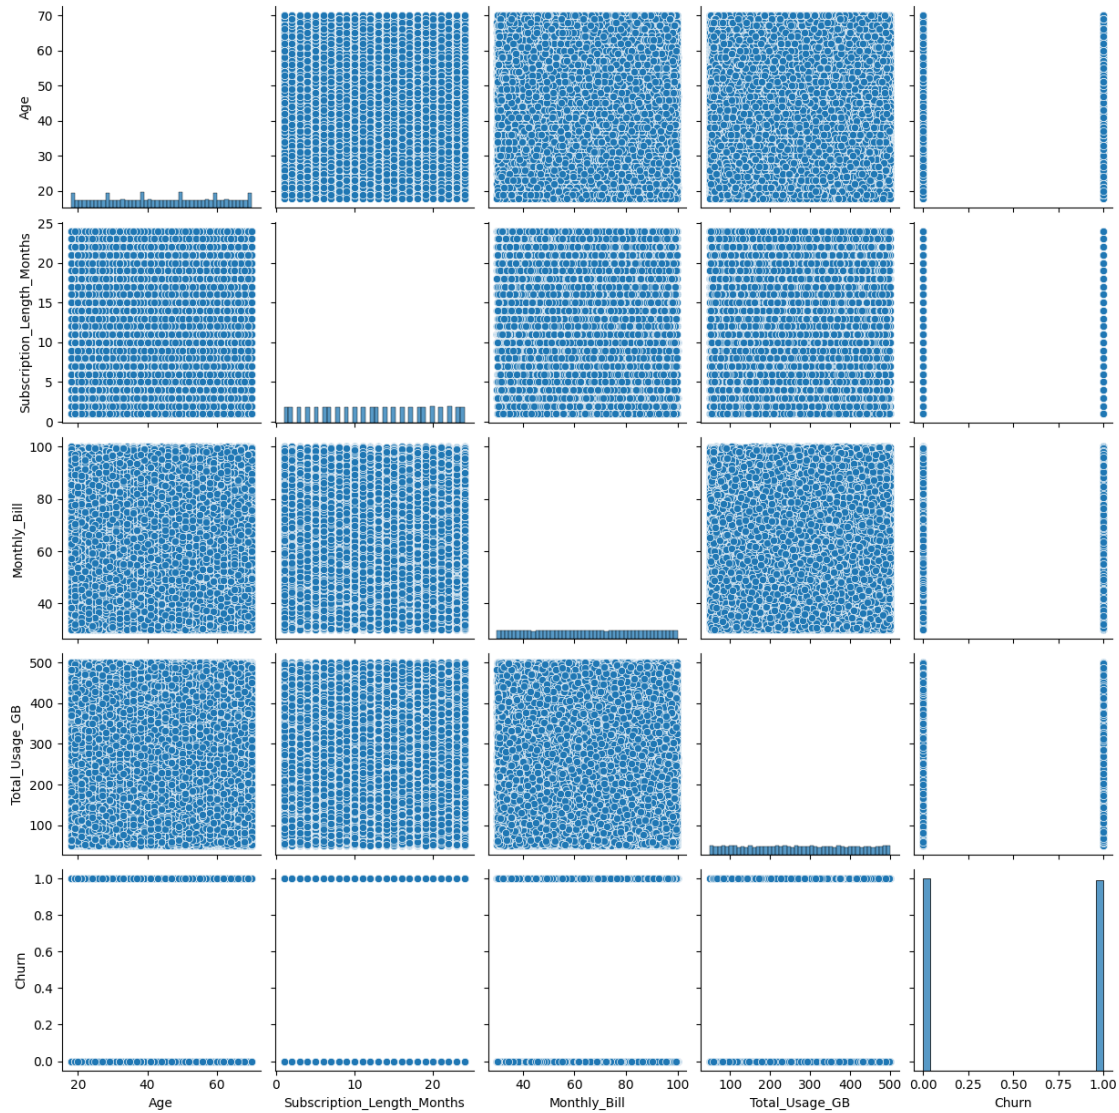
# Remove empty subplots
for i in range(len(numerical_columns), num_rows * num_plots_per_row):
    fig.delaxes(axes.flatten()[i])

plt.tight_layout()
plt.show()

```



```
[ ]: # Data visualization (e.g., histograms, scatter plots)
sns.pairplot(df)
plt.show()
```



No missing values found

```
[ ]: df.isnull().sum()
```

```
[ ]: Age          0
      Gender       0
      Location     0
      Subscription_Length_Months  0
      Monthly_Bill  0
      Total_Usage_GB  0
      Churn        0
      Senior Citizen  0
      Junior Citizen  0
```



```
Median Citizen          0
dtype: int64
```

```
[ ]: df
```

```
[ ]:
      CustomerID      Name  Age  Gender  Location \
0              1  Customer_1  63   Male  Los Angeles
1              2  Customer_2  62  Female    New York
2              3  Customer_3  24  Female  Los Angeles
3              4  Customer_4  36  Female    Miami
4              5  Customer_5  46  Female    Miami
...
99995      99996  Customer_99996  33   Male    Houston
99996      99997  Customer_99997  62  Female    New York
99997      99998  Customer_99998  64   Male    Chicago
99998      99999  Customer_99999  51  Female    New York
99999     100000  Customer_100000  27  Female  Los Angeles

      Subscription_Length_Months  Monthly_Bill  Total_Usage_GB  Churn  \
0                               17          73.36           236      0
1                               1          48.76           172      0
2                               5          85.47           460      0
3                               3          97.94           297      1
4                               19          58.14           266      0
...
99995                        23          55.13           226      1
99996                        19          61.65           351      0
99997                        17          96.11           251      1
99998                        20          49.25           434      1
99999                        19          76.57           173      1

      Senior Citizen  Junior Citizen  Median Citizen  SubscriptionCategory
0                  1                0                0          Two Years
1                  1                0                0          One Month
2                  0                0                1          One Year
3                  0                0                1          One Year
4                  0                0                1          Two Years
...
99995              0                0                1          Two Years
99996              1                0                0          Two Years
99997              1                0                0          Two Years
99998              1                0                0          Two Years
99999              0                0                1          Two Years
```

```
[100000 rows x 13 columns]
```

```
[ ]: df_encoded.shape
```

```
[ ]: (100000, 15)
```

```
[ ]: df_encoded
```

```
[ ]:      Age  Subscription_Length_Months  Monthly_Bill  Total_Usage_GB  Churn  \
0      63                             17          73.36           236      0
1      62                             1          48.76           172      0
2      24                             5          85.47           460      0
3      36                             3          97.94           297      1
4      46                             19          58.14           266      0
...    ...
99995   33                             23          55.13           226      1
99996   62                             19          61.65           351      0
99997   64                             17          96.11           251      1
99998   51                             20          49.25           434      1
99999   27                             19          76.57           173      1
```

```
      Senior Citizen  Junior Citizen  Median Citizen  Gender_Female  \
0                  1                0                0                0
1                  1                0                0                1
2                  0                0                1                1
3                  0                0                1                1
4                  0                0                1                1
...
99995              ...              ...              ...              ...
99995              0                0                1                0
99996              1                0                0                1
99997              1                0                0                0
99998              1                0                0                1
99999              0                0                1                1
```

```
      Gender_Male  Location_Chicago  Location_Houston  Location_Los Angeles  \
0                1                  0                  0                  1
1                0                  0                  0                  0
2                0                  0                  0                  1
3                0                  0                  0                  0
4                0                  0                  0                  0
...
99995              ...              ...              ...              ...
99995              1                  0                  1                  0
99996              0                  0                  0                  0
99997              1                  1                  0                  0
99998              0                  0                  0                  0
99999              0                  0                  0                  1
```

```
      Location_Miami  Location_New York
0                  0                  0
1                  0                  1
2                  0                  0
```

3	1	0
4	1	0
...
99995	0	0
99996	0	1
99997	0	0
99998	0	1
99999	0	0

[100000 rows x 15 columns]

Outlier Removal is not needed yet but can be done with z score method by uncommenting this:

```
[ ]: # z_scores.describe()
# # Example: Remove outliers using Z-score
# z_scores = zscore(df_encoded)
# df_no_outliers = df_encoded[(abs(z_scores) < 3).all(axis=1)]
# df_no_outliers
```

We could say looking at the above result that this doesn't have outliers

Normalisation did not have a lot of effect here on the dataset

We will now do encoding of categorical variables in our data

1.0.1 Encoding

```
[ ]: from sklearn.preprocessing import LabelEncoder

# Create a LabelEncoder instance
label_encoder = LabelEncoder()
# Apply label encoding to the 'Age_Group_Indicator' column
df['Age_Group_Indicator'] = label_encoder.
    ↪fit_transform(df['Age_Group_Indicator'])
df = pd.get_dummies(df, columns=['Gender', 'Location'])
```

```
[ ]: df.head()
```

	CustomerID	Name	Age	Subscription_Length_Months	Monthly_Bill	\
0	1	Customer_1	63	17	73.36	
1	2	Customer_2	62	1	48.76	
2	3	Customer_3	24	5	85.47	
3	4	Customer_4	36	3	97.94	
4	5	Customer_5	46	19	58.14	

	Total_Usage_GB	Churn	Average_Monthly_Data_Usage	Billing_Change_Rate	\
0	236	0	13.882353	NaN	
1	172	0	172.000000	-24.60	

2	460	0	92.000000	36.71
3	297	1	99.000000	12.47
4	266	0	14.000000	-39.80

	Billing_As_Percentage	Customer_Tenure_Months	Churn_History	\
0	112.769247	17	NaN	
1	74.954041	1	0.0	
2	131.384781	5	0.0	
3	150.553708	3	0.0	
4	89.373010	19	1.0	

	Age_Group_Indicator	Remaining_Subscription_Length	Average_Bill_Change	\
0	1	17	NaN	
1	1	0	NaN	
2	2	3	NaN	
3	0	0	8.193333	
4	0	15	3.126667	

	Gender_Female	Gender_Male	Location_Chicago	Location_Houston	\
0	0	1	0	0	
1	1	0	0	0	
2	1	0	0	0	
3	1	0	0	0	
4	1	0	0	0	

	Location_Los Angeles	Location_Miami	Location_New York
0	1	0	0
1	0	0	1
2	1	0	0
3	0	1	0
4	0	1	0

```
[ ]: df.describe()
```

```
[ ]:
```

	CustomerID	Age	Subscription_Length_Months	\
count	100000.000000	100000.000000	100000.000000	
mean	50000.500000	44.027020	12.490100	
std	28867.657797	15.280283	6.926461	
min	1.000000	18.000000	1.000000	
25%	25000.750000	31.000000	6.000000	
50%	50000.500000	44.000000	12.000000	
75%	75000.250000	57.000000	19.000000	
max	100000.000000	70.000000	24.000000	

	Monthly_Bill	Total_Usage_GB	Churn	\
count	100000.000000	100000.000000	100000.000000	
mean	65.053197	274.393650	0.497790	

std	20.230696	130.463063	0.499998
min	30.000000	50.000000	0.000000
25%	47.540000	161.000000	0.000000
50%	65.010000	274.000000	0.000000
75%	82.640000	387.000000	1.000000
max	100.000000	500.000000	1.000000

	Average_Monthly_Data_Usage	Billing_Change_Rate	Billing_As_Percentage \
count	100000.000000	99999.000000	100000.000000
mean	43.349682	0.000032	100.000000
std	65.786483	28.588777	31.098696
min	2.083333	-69.910000	46.116104
25%	12.687500	-20.680000	73.078653
50%	21.909091	-0.030000	99.933598
75%	42.545455	20.580000	127.034495
max	500.000000	69.460000	153.720347

	Customer_Tenure_Months	Churn_History	Age_Group_Indicator \
count	100000.000000	99999.000000	100000.000000
mean	12.490100	0.497785	0.866270
std	6.926461	0.499998	0.777375
min	1.000000	0.000000	0.000000
25%	6.000000	0.000000	0.000000
50%	12.000000	0.000000	1.000000
75%	19.000000	1.000000	1.000000
max	24.000000	1.000000	2.000000

	Remaining_Subscription_Length	Average_Bill_Change	Gender_Female \
count	100000.000000	99997.000000	100000.000000
mean	-49987.009900	0.000048	0.502160
std	28867.620918	9.551407	0.499998
min	-99986.000000	-23.286667	0.000000
25%	-74986.250000	-6.856667	0.000000
50%	-49988.000000	-0.020000	1.000000
75%	-24985.250000	6.876667	1.000000
max	17.000000	23.166667	1.000000

	Gender_Male	Location_Chicago	Location_Houston \
count	100000.000000	100000.000000	100000.000000
mean	0.497840	0.199580	0.201570
std	0.499998	0.399687	0.401175
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	1.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000

	Location_Los Angeles	Location_Miami	Location_New York
count	100000.000000	100000.000000	100000.000000
mean	0.200410	0.200310	0.198130
std	0.400309	0.400234	0.398593
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000

```
[ ]: from sklearn.model_selection import train_test_split

# Define your target variable and features
X = df.drop('Churn', axis=1)
y = df['Churn']
```

```
[ ]: from sklearn.preprocessing import MinMaxScaler

# # # Create a MinMaxScaler object
scaler = MinMaxScaler()

# # # Fit the scaler to your data and transform it
X_normalized = scaler.fit_transform(X)
```

Various methods with/without normalisation and with/without extra features were tried and tested here

```
[ ]: # Without Normalisation
X_train2, X_test2, y_train2, y_test2 = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

```
[ ]: imputer = SimpleImputer(strategy='mean')
X_imputed = imputer.fit_transform(X)
X_train3, X_test3, y_train3, y_test3 = train_test_split(X_imputed, y,
↳test_size=0.2, random_state=42)
```

```
[ ]: # Split the data into training and testing sets (e.g., 80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X_normalized, y,
↳test_size=0.2, random_state=42)
```

```
[ ]: X_train[0]
```

```
[ ]: array([0.69230769, 0.17391304, 0.77857143, 0.34444444, 1.
,
0.
, 0.
, 1.
, 0.
, 0.
,
0.
, 0.
, 0.
, 1.
])
```

2 2 Feature Engineering:

- Generate relevant features from the dataset that can help improve the model's prediction accuracy.
- Apply feature scaling or normalization if necessary.

```
[ ]: import pandas as pd

# Assuming you have a DataFrame 'df' with the original features: Age,
#   ↳ Subscription_Length_Months, Monthly_Bill, Total_Usage_GB
# Replace 'your_dataset.csv' with the path to your dataset
# df = pd.read_csv('your_dataset.csv')

# Feature Engineering

# 1. Average Monthly Data Usage
df['Average_Monthly_Data_Usage'] = df['Total_Usage_GB'] /
#   ↳ df['Subscription_Length_Months']

# 2. Billing Change Rate
df['Billing_Change_Rate'] = df['Monthly_Bill'].diff()

# 3. Billing Amount as a Percentage
df['Billing_As_Percentage'] = (df['Monthly_Bill'] / df['Monthly_Bill'].mean())
#   ↳ * 100

# 4. Customer Tenure in Months
df['Customer_Tenure_Months'] = df['Subscription_Length_Months']

# 5. Churn History (Assuming 'Churn' is a binary column indicating churn
#   ↳ history)
df['Churn_History'] = df['Churn'].shift(1) # Lagged version of the churn column

# 6. Age Group Indicator (Assuming age groups are defined)
age_bins = [0, 30, 50, 100] # Define your age groups as needed
age_labels = ['Young', 'Middle-Aged', 'Senior']
df['Age_Group_Indicator'] = pd.cut(df['Age'], bins=age_bins, labels=age_labels)

# 7. Remaining Subscription Length
df['Remaining_Subscription_Length'] = df['Subscription_Length_Months'] - df.
#   ↳ index

# 8. Average Bill Change
df['Average_Bill_Change'] = df['Billing_Change_Rate'].rolling(window=3).mean()

# Display the updated DataFrame with engineered features
print(df.head())
```

	CustomerID	Name	Age	Gender	Location	\
0	1	Customer_1	63	Male	Los Angeles	
1	2	Customer_2	62	Female	New York	
2	3	Customer_3	24	Female	Los Angeles	
3	4	Customer_4	36	Female	Miami	
4	5	Customer_5	46	Female	Miami	

	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	Churn	\
0	17	73.36	236	0	
1	1	48.76	172	0	
2	5	85.47	460	0	
3	3	97.94	297	1	
4	19	58.14	266	0	

	Average_Monthly_Data_Usage	Billing_Change_Rate	Billing_As_Percentage	\
0	13.882353	NaN	112.769247	
1	172.000000	-24.60	74.954041	
2	92.000000	36.71	131.384781	
3	99.000000	12.47	150.553708	
4	14.000000	-39.80	89.373010	

	Customer_Tenure_Months	Churn_History	Age_Group_Indicator	\
0	17	NaN	Senior	
1	1	0.0	Senior	
2	5	0.0	Young	
3	3	0.0	Middle-Aged	
4	19	1.0	Middle-Aged	

	Remaining_Subscription_Length	Average_Bill_Change
0	17	NaN
1	0	NaN
2	3	NaN
3	0	8.193333
4	15	3.126667

```
[ ]: df.drop(columns=['Name', 'CustomerID'], axis=1, inplace=True)
```

```
[ ]: df.columns
```

```
[ ]: Index(['Age', 'Gender', 'Location', 'Subscription_Length_Months',
           'Monthly_Bill', 'Total_Usage_GB', 'Churn', 'Average_Monthly_Data_Usage',
           'Billing_Change_Rate', 'Billing_As_Percentage',
           'Customer_Tenure_Months', 'Churn_History', 'Age_Group_Indicator',
           'Remaining_Subscription_Length', 'Average_Bill_Change'],
          dtype='object')
```


2.0.1 Graphical Analysis

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load your dataset
# df = pd.read_csv('customer_data.csv') # Replace 'customer_data.csv' with
↳ your dataset file path

# Assuming your dataset has columns like 'age', 'churn', 'city', and 'gender'

# Filter the data by city and gender
cities = df['Location'].unique()
genders = df['Gender'].unique()

# Create subplots for different cities and genders
fig, axes = plt.subplots(len(cities), len(genders), figsize=(15, 10),
↳ sharex=True, sharey=True)

for i, city in enumerate(cities):
    for j, gender in enumerate(genders):
        # Filter data for the current city and gender
        subset = df[(df['Location'] == city) & (df['Gender'] == gender)]

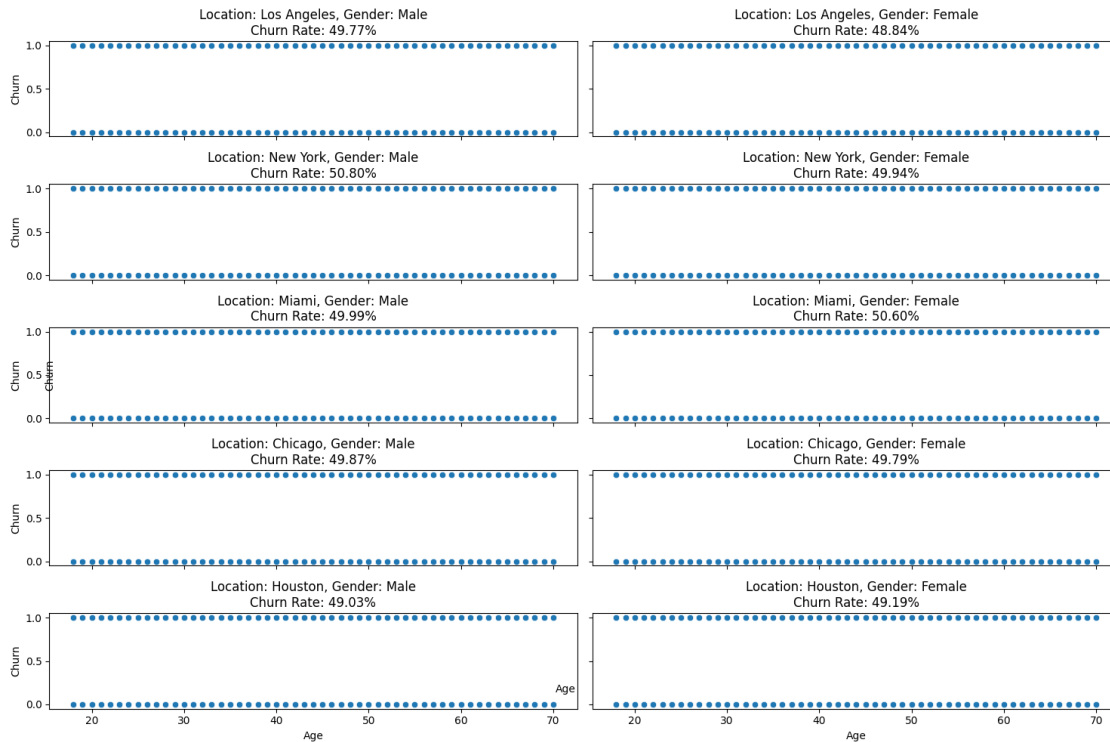
        # Calculate churn rate for the subset
        churn_rate = subset['Churn'].mean()

        # Create a scatter plot
        sns.scatterplot(data=subset, x='Age', y='Churn', ax=axes[i, j])
        axes[i, j].set_title(f'Location: {city}, Gender: {gender}\nChurn Rate:
↳ {churn_rate:.2%}')

# Add labels and a common y-axis label
fig.text(0.5, 0.08, 'Age', ha='center')
fig.text(0.04, 0.5, 'Churn', va='center', rotation='vertical')

# Adjust subplot spacing
plt.tight_layout()

# Show the plot
plt.show()
```



```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

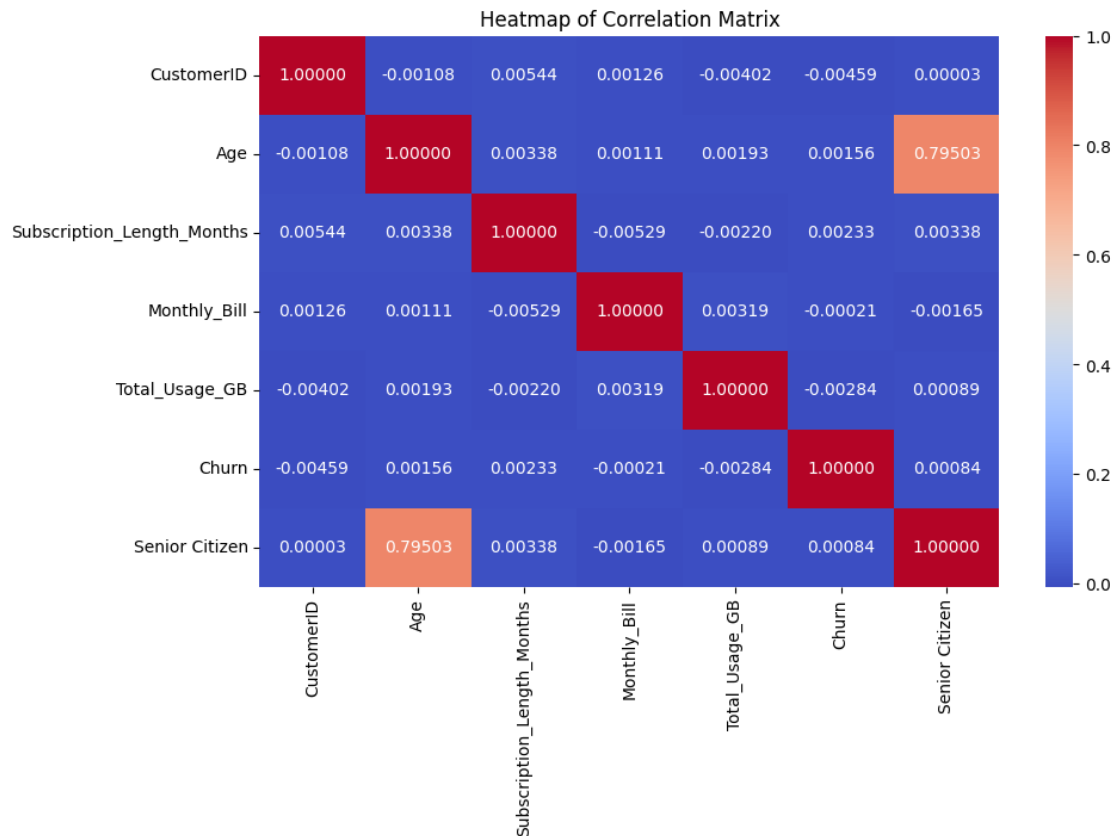
# Load your dataset
# df = pd.read_csv('customer_data.csv') # Replace 'customer_data.csv' with
↳ your dataset file path

# Assuming your dataset has columns like 'age', 'churn', 'city', and 'gender'

# Define the age threshold for senior citizens
senior_age_threshold = 55

# Create a new column 'Senior Citizen' based on age
df['Senior Citizen'] = df['Age'] >= senior_age_threshold

# Create a heatmap to visualize the relationship between 'Senior Citizen' and
↳ other variables
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".5f")
plt.title('Heatmap of Correlation Matrix')
plt.show()
```



```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load your dataset
# df = pd.read_csv('customer_data.csv') # Replace 'customer_data.csv' with
# your dataset file path

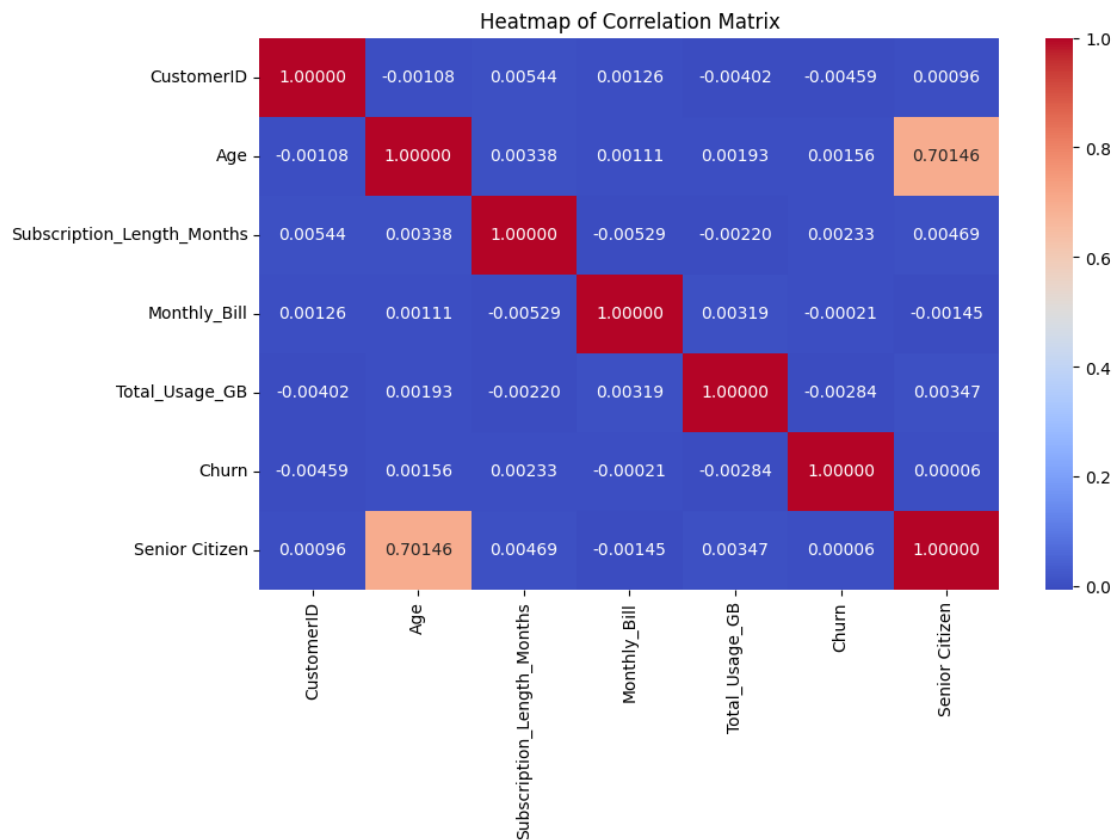
# Assuming your dataset has columns like 'age', 'churn', 'city', and 'gender'

# Define the age threshold for senior citizens
senior_age_threshold = 60

# Create a new column 'Senior Citizen' based on age
df['Senior Citizen'] = df['Age'] >= senior_age_threshold

# Create a heatmap to visualize the relationship between 'Senior Citizen' and
# other variables
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".5f")
```

```
plt.title('Heatmap of Correlation Matrix')
plt.show()
```



```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

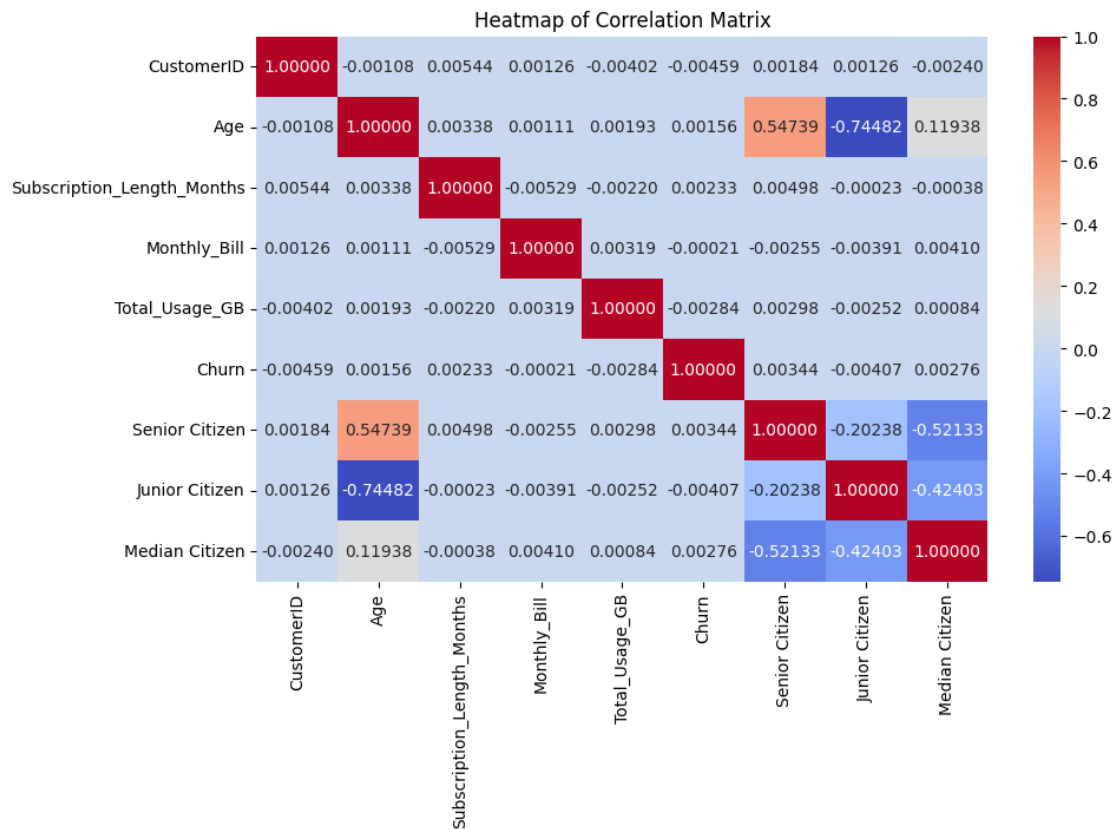
# Load your dataset
# df = pd.read_csv('customer_data.csv') # Replace 'customer_data.csv' with
# your dataset file path

# Assuming your dataset has columns like 'age', 'churn', 'city', and 'gender'

# Define the age threshold for senior citizens
senior_age_threshold = 65

# Create a new column 'Senior Citizen' based on age
df['Senior Citizen'] = df['Age'] >= senior_age_threshold
```

```
# Create a heatmap to visualize the relationship between 'Senior Citizen' and
↳ other variables
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".5f")
plt.title('Heatmap of Correlation Matrix')
plt.show()
```



```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load your dataset
# df = pd.read_csv('customer_data.csv') # Replace 'customer_data.csv' with
↳ your dataset file path

# Assuming your dataset has columns like 'age', 'churn', 'city', and 'gender'

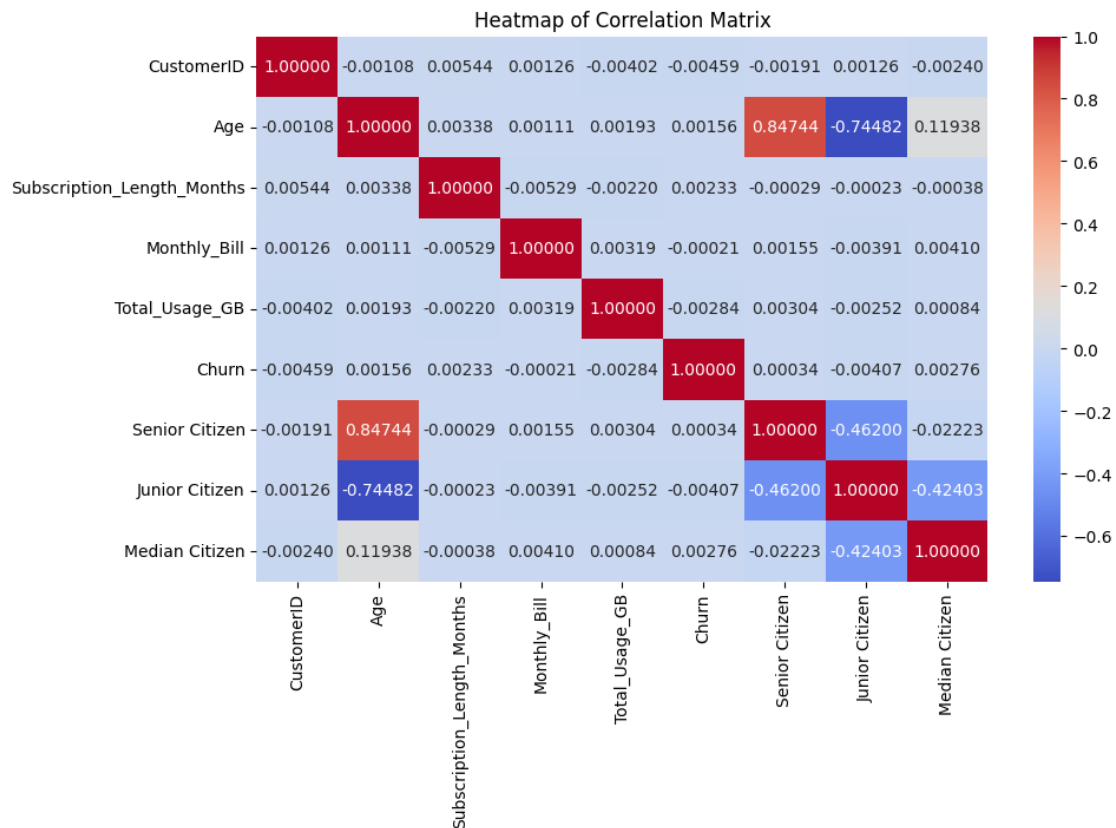
# Define the age threshold for senior citizens
senior_age_threshold = 50
```

```

# Create a new column 'Senior Citizen' based on age
df['Senior Citizen'] = df['Age'] >= senior_age_threshold

# Create a heatmap to visualize the relationship between 'Senior Citizen' and
↳ other variables
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".5f")
plt.title('Heatmap of Correlation Matrix')
plt.show()

```



```

[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load your dataset
# df = pd.read_csv('customer_data.csv') # Replace 'customer_data.csv' with
↳ your dataset file path

# Assuming your dataset has columns like 'age', 'churn', 'city', and 'gender'

```


```

# Define the age threshold for senior citizens
junior_age_threshold = 25

# Create a new column 'Senior Citizen' based on age
df['Junior Citizen'] = df['Age'] <= junior_age_threshold

# Create a heatmap to visualize the relationship between 'Senior Citizen' and
↳ other variables
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".5f")
plt.title('Heatmap of Correlation Matrix')
plt.show()

```



```

[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load your dataset
# df = pd.read_csv('customer_data.csv') # Replace 'customer_data.csv' with
↳ your dataset file path

```

```


# Assuming your dataset has columns like 'age', 'churn', 'city', and 'gender'

# Define the age threshold for senior citizens
junior_age_threshold = 25

# Create a new column 'Senior Citizen' based on age
df['Junior Citizen'] = df['Age'] <= junior_age_threshold

# Create a heatmap to visualize the relationship between 'Senior Citizen' and
↳ other variables
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".5f")
plt.title('Heatmap of Correlation Matrix')
plt.show()

```



```

[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load your dataset

```




```
# df = pd.read_csv('customer_data.csv') # Replace 'customer_data.csv' with
↳ your dataset file path

# Assuming your dataset has columns like 'age', 'churn', 'city', and 'gender'

# Define the age threshold for senior citizens
junior_age_threshold = 30

# Create a new column 'Senior Citizen' based on age
df['Junior Citizen'] = df['Age'] <= junior_age_threshold

# Create a heatmap to visualize the relationship between 'Senior Citizen' and
↳ other variables
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".5f")
plt.title('Heatmap of Correlation Matrix')
plt.show()
```



```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
```

```

import seaborn as sns

# Load your dataset
# df = pd.read_csv('customer_data.csv') # Replace 'customer_data.csv' with
    ↳ your dataset file path

# Assuming your dataset has columns like 'age', 'churn', 'city', and 'gender'

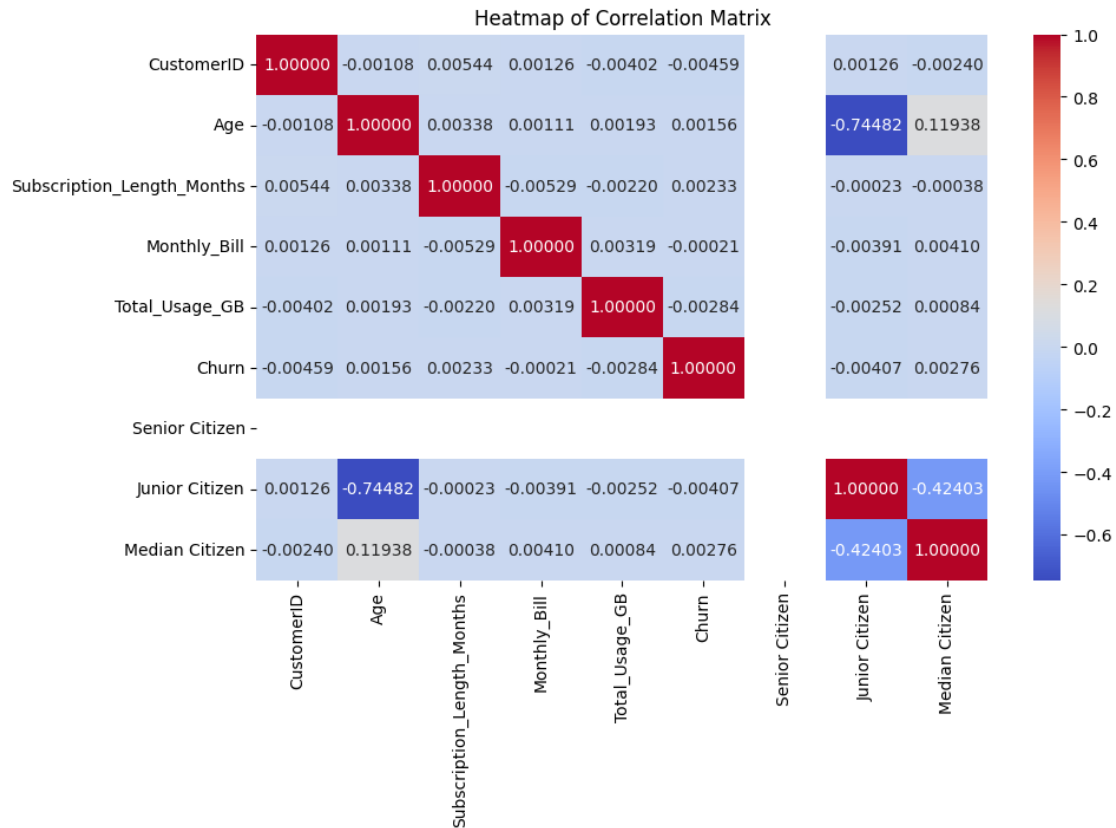
# Define the age threshold for senior citizens
high_age_threshold = 65
low_age_threshold = 25

# Create a new column 'Senior Citizen' based on age
# df['Median Citizen'] = df[low_age_threshold <= df['Age'] <=
    ↳ high_age_threshold]

df['Median Citizen'] = df['Age'].apply(lambda age: high_age_threshold >=age >=
    ↳ low_age_threshold)


# Create a heatmap to visualize the relationship between 'Senior Citizen' and
    ↳ other variables
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".5f")
plt.title('Heatmap of Correlation Matrix')
plt.show()

```



```
[ ]: correlation_matrix = pd.DataFrame(df,columns=df.columns).corr()

# Create a heatmap
plt.figure(figsize=(15, 12)) # Adjust the figure size as needed
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".5f")
plt.title('Correlation Heatmap')
plt.show()
```



Surprisingly we see we do not have multicollinearity amongst our features

more usage => more bill

more usage => less churn

more bill => less churn

age => increase in churn

```
[ ]: # df['Senior Citizen'] = df['Age'] >= 50
      # df['Junior Citizen'] = df['Age'] <= 20
      # df['Median Citizen'] = df['Age'].apply(lambda age: 50 >= age >= 20)

      ## Convert boolean columns to integers (0 or 1)
      # df['Senior Citizen'] = df['Senior Citizen'].astype(int)
      # df['Junior Citizen'] = df['Junior Citizen'].astype(int)
      # df['Median Citizen'] = df['Median Citizen'].astype(int)
```

```
# df
```

```
[ ]:      CustomerID      Name  Age  Gender  Location \
0          1  Customer_1  63    Male  Los Angeles
1          2  Customer_2  62   Female   New York
2          3  Customer_3  24   Female  Los Angeles
3          4  Customer_4  36   Female    Miami
4          5  Customer_5  46   Female    Miami
...
99995      99996  Customer_99996  33    Male    Houston
99996      99997  Customer_99997  62   Female   New York
99997      99998  Customer_99998  64    Male    Chicago
99998      99999  Customer_99999  51   Female   New York
99999     100000  Customer_100000  27   Female  Los Angeles

      Subscription_Length_Months  Monthly_Bill  Total_Usage_GB  Churn \
0                               17          73.36           236      0
1                               1          48.76           172      0
2                               5          85.47           460      0
3                               3          97.94           297      1
4                               19          58.14           266      0
...
99995                        23          55.13           226      1
99996                        19          61.65           351      0
99997                        17          96.11           251      1
99998                        20          49.25           434      1
99999                        19          76.57           173      1

      Senior Citizen  Junior Citizen  Median Citizen
0                   1                0                0
1                   1                0                0
2                   0                0                1
3                   0                0                1
4                   0                0                1
...
99995                0                0                1
99996                1                0                0
99997                1                0                0
99998                1                0                0
99999                0                0                1
```

```
[100000 rows x 12 columns]
```

```
[ ]: colors = ['#4D3425', '#E4512B']
seniority_churn = df.groupby(['Senior Citizen', 'Churn']).size().unstack()

ax = (seniority_churn.T*100.0 / seniority_churn.T.sum()).T.plot(kind='bar',
```

```

width = 0.2,
stacked = True,
rot = 0,
figsize = (8,6),
color = colors)

ax.yaxis.set_major_formatter(mtick.PercentFormatter())
ax.legend(loc='center',prop={'size':14},title = 'Churn')
ax.set_ylabel('% Customers')
ax.set_title('Churn by Seniority Level',size = 14)

# Code to add the data labels on the stacked bar chart
for p in ax.patches:
    width, height = p.get_width(), p.get_height()
    x, y = p.get_xy()
    ax.annotate('{:.0f}%'.format(height), (p.get_x()+.25*width, p.get_y()+.
↪4*height),
                color = 'white',
                weight = 'bold',size =14)


```

```

-----
NameError                                Traceback (most recent call last)
<ipython-input-72-18b65e94b1fe> in <cell line: 10>()
      8                                     figsize,
↪ = (8,6),
      9                                     color =
↪ colors)
--> 10 ax.yaxis.set_major_formatter(mtick.PercentFormatter())
     11 ax.legend(loc='center',prop={'size':14},title = 'Churn')
     12 ax.set_ylabel('% Customers')

NameError: name 'mtick' is not defined

```



```
[ ]: colors = ['#4D3425', '#E4512B']
seniority_churn = df.groupby(['Junior Citizen', 'Churn']).size().unstack()

ax = (seniority_churn.T*100.0 / seniority_churn.T.sum()).T.plot(kind='bar',
                                                                width = 0.2,
                                                                stacked = True,
                                                                rot = 0,
                                                                figsize = (8,6),
                                                                color = colors)

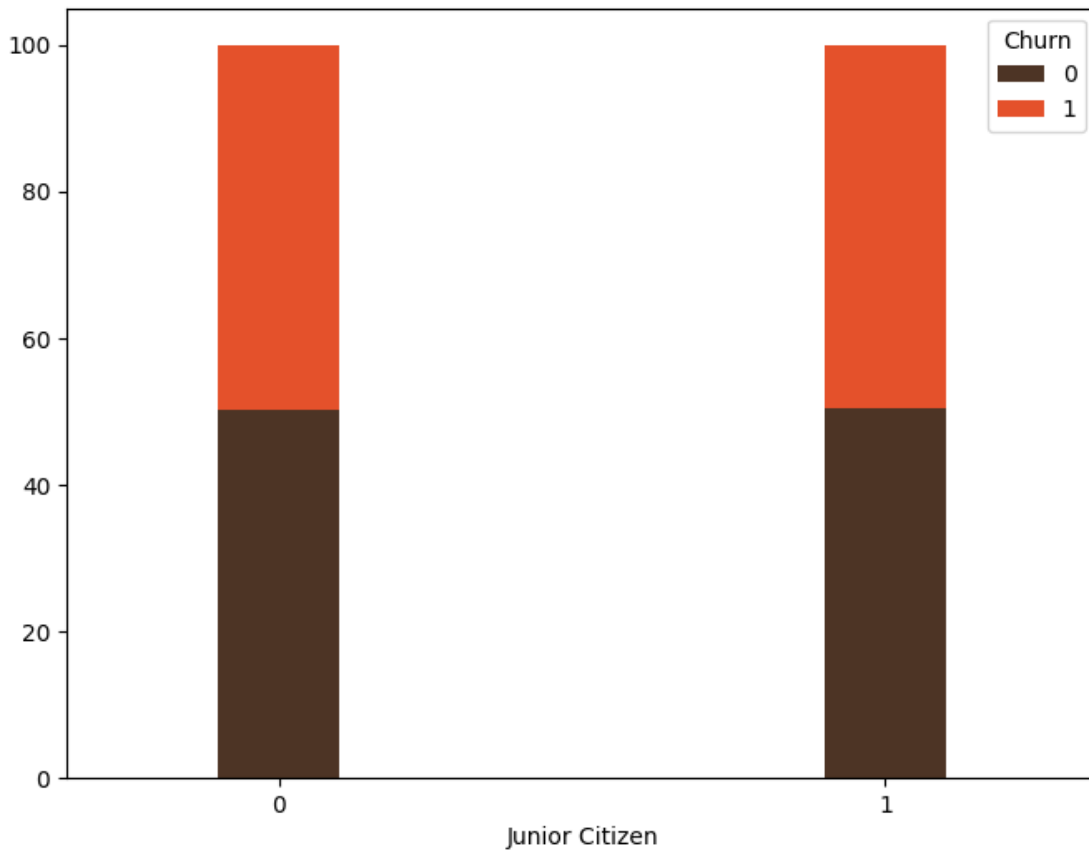
ax.yaxis.set_major_formatter(mtick.PercentFormatter())
ax.legend(loc='center',prop={'size':14},title = 'Churn')
ax.set_ylabel('% Customers')
ax.set_title('Churn by Seniority Level',size = 14)

# Code to add the data labels on the stacked bar chart
for p in ax.patches:
    width, height = p.get_width(), p.get_height()
    x, y = p.get_xy()
    ax.annotate('{:.0f}%'.format(height), (p.get_x()+.25*width, p.get_y()+.
↪4*height),
```

```
color = 'white',
weight = 'bold',size =14)
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-73-f80bf588d86d> in <cell line: 10>()
      8                                     figsize=(8,6),
    => 9                                     color = colors)
    --> 10 ax.yaxis.set_major_formatter(mtick.PercentFormatter())
      11 ax.legend(loc='center',prop={'size':14},title = 'Churn')
      12 ax.set_ylabel('% Customers')

NameError: name 'mtick' is not defined
```



```
[ ]: colors = ['#4D3425','#E4512B']
seniority_churn = df.groupby(['Median Citizen','Churn']).size().unstack()
```



```

ax = (seniority_churn.T*100.0 / seniority_churn.T.sum()).T.plot(kind='bar',
                                                                width = 0.2,
                                                                stacked = True,
                                                                rot = 0,
                                                                figsize = (8,6),
                                                                color = colors)

ax.yaxis.set_major_formatter(mtick.PercentFormatter())
ax.legend(loc='center',prop={'size':14},title = 'Churn')
ax.set_ylabel('% Customers')
ax.set_title('Churn by Seniority Level',size = 14)

# Code to add the data labels on the stacked bar chart
for p in ax.patches:
    width, height = p.get_width(), p.get_height()
    x, y = p.get_xy()
    ax.annotate('{:.0f}%'.format(height), (p.get_x()+.25*width, p.get_y()+
↪4*height),
                color = 'white',
                weight = 'bold',size =14)

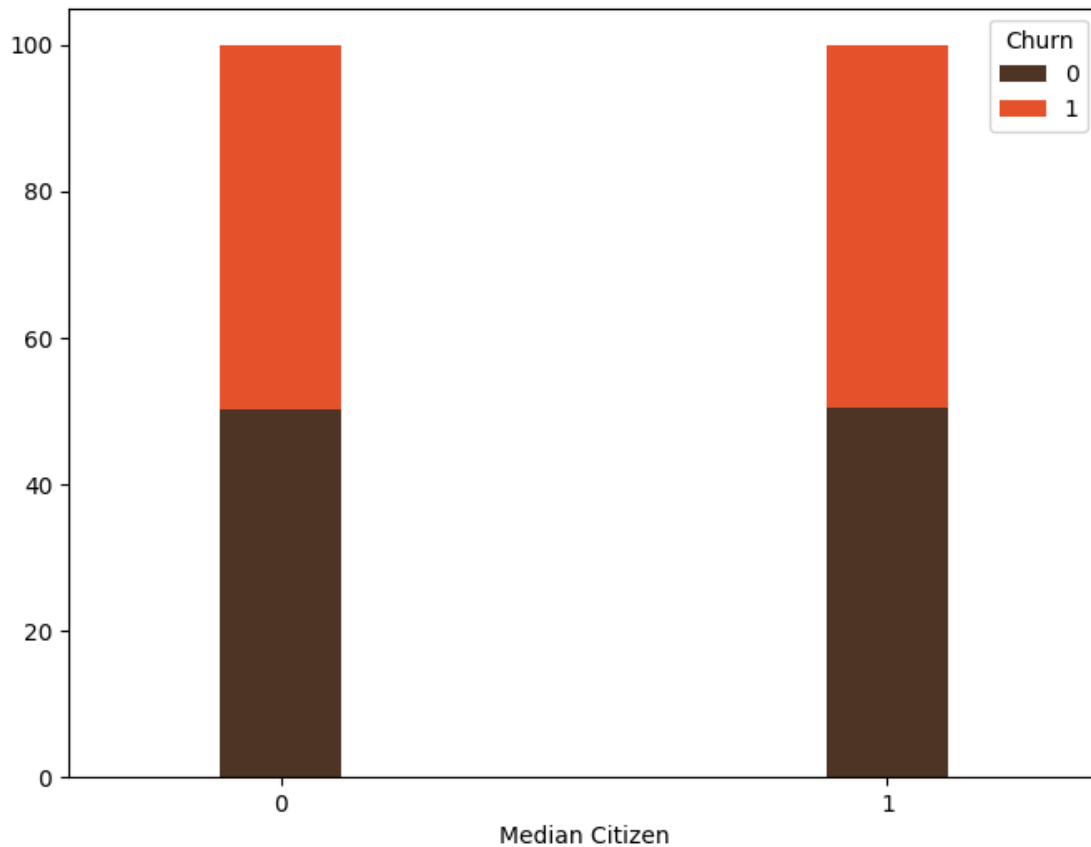
```

```

-----
NameError                                Traceback (most recent call last)
<ipython-input-74-dfa4f5ff82e9> in <cell line: 10>()
      8                                figsize,
↪ = (8,6),
      9                                color =
↪ colors)
----> 10 ax.yaxis.set_major_formatter(mtick.PercentFormatter())
      11 ax.legend(loc='center',prop={'size':14},title = 'Churn')
      12 ax.set_ylabel('% Customers')

NameError: name 'mtick' is not defined

```



```
[ ]: # Create a new column based on the duration
```

```
def categorize_subscription(duration):
    if duration <= 1:
        return 'One Month'
    elif duration <= 12:
        return 'One Year'
    elif duration <= 24:
        return 'Two Years'
    else:
        return 'Greater than Two Years'
```

```
df['SubscriptionCategory'] = df['Subscription_Length_Months'].
    ↪ apply(categorize_subscription)
```

```
[ ]: colors = ['#4D3425', '#E4512B']
```

```
seniority_churn = df.groupby(['SubscriptionCategory', 'Churn']).size().unstack()
```

```
ax = (seniority_churn.T*100.0 / seniority_churn.T.sum()).T.plot(kind='bar',
                                                                width = 0.2,
                                                                stacked = True,
```

```

rot = 0,
figsize = (8,6),
color = colors)

ax.yaxis.set_major_formatter(mtick.PercentFormatter())
ax.legend(loc='center',prop={'size':14},title = 'Churn')
ax.set_ylabel('% Customers')
ax.set_title('Churn by Seniority Level',size = 14)

# Code to add the data labels on the stacked bar chart
for p in ax.patches:
    width, height = p.get_width(), p.get_height()
    x, y = p.get_xy()
    ax.annotate('{:.0f}%'.format(height), (p.get_x()+.25*width, p.get_y()+.
↪4*height),
                color = 'white',
                weight = 'bold',size =14)


```

```

-----
NameError                                Traceback (most recent call last)
<ipython-input-78-8aee91c6acb7> in <cell line: 10>()
      8                                     figsize
↪ = (8,6),
      9                                     color =
↪ colors)
----> 10 ax.yaxis.set_major_formatter(mtick.PercentFormatter())
      11 ax.legend(loc='center',prop={'size':14},title = 'Churn')
      12 ax.set_ylabel('% Customers')

NameError: name 'mtick' is not defined

```

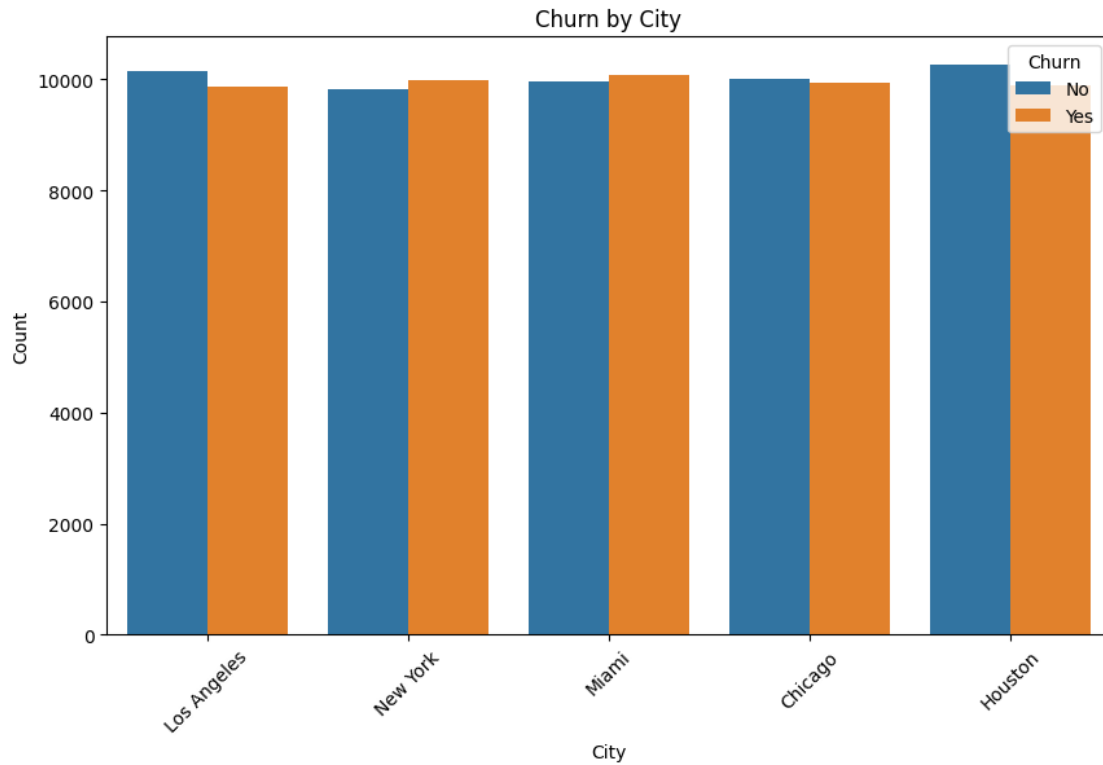


```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load your dataset
# df = pd.read_csv('customer_data.csv') # Replace 'customer_data.csv' with
# your dataset file path

# Assuming your dataset includes 'Churn' and 'City' columns



# Create a countplot to visualize the relationship between cities and churn
plt.figure(figsize=(10, 6))
sns.countplot(x='Location', hue='Churn', data=df)
plt.title('Churn by City')
plt.xlabel('City')
plt.ylabel('Count')
plt.xticks(rotation=45) # Rotate x-axis labels for better readability if needed
plt.legend(title='Churn', labels=['No', 'Yes'])
plt.show()
```





```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

continuous=['Age', 'Subscription_Length_Months', 'Monthly_Bill', '
↳ 'Total_Usage_GB']
categorical=[ 'Gender', 'Location']

for col in continuous:
    plt.figure(figsize=(8, 5))
    ax = sns.kdeplot(df[col][df["Churn"] == 0], color="Red", shade=True)
    ax = sns.kdeplot(df[col][df["Churn"] == 1], ax=ax, color="Blue", shade=True)
    ax.legend(["Not Churn", "Churn"], loc='upper right')
    ax.set_ylabel('Density')
    ax.set_xlabel(col)
    ax.set_title(f'Distribution of {col} by Churn')
    plt.show()
```





```
[ ]: # todo
```

```
[ ]: df
```

```
[ ]:      Age  Gender  Location  Subscription_Length_Months  Monthly_Bill  \
0      63   Male  Los Angeles              17             73.36
1      62  Female   New York               1             48.76
2      24  Female  Los Angeles               5             85.47
3      36  Female    Miami                3             97.94
4      46  Female    Miami               19             58.14
...    ...    ...    ...
99995   33   Male    Houston              23             55.13
99996   62  Female   New York              19             61.65
99997   64   Male   Chicago              17             96.11
99998   51  Female   New York              20             49.25
99999   27  Female  Los Angeles              19             76.57
```

```
      Total_Usage_GB  Churn  Senior Citizen  Junior Citizen  Median Citizen
0                236      0                1                0                0
1                172      0                1                0                0
2                460      0                0                0                1
3                297      1                0                0                1
4                266      0                0                0                1
...    ...    ...    ...    ...    ...
99995            226      1                0                0                1
99996            351      0                1                0                0
99997            251      1                1                0                0
99998            434      1                1                0                0
99999            173      1                0                0                1
```

```
[100000 rows x 10 columns]
```

```
[ ]: colors = ['#E94B3C','#2D2926']


churn = df[df['Churn'] == 1].describe().T
not_churn = df[df['Churn'] == 0].describe().T

fig,ax = plt.subplots(nrows = 1,ncols = 2,figsize = (8,8))
plt.subplot(1,2,1)
sns.heatmap(churn[['mean']],annot = True,cmap = colors,linewidths = 0.
    ↪4,linecolor = 'black',cbar = False,fmt = '.2f')
plt.title('Churned Customers');

plt.subplot(1,2,2)
sns.heatmap(not_churn[['mean']],annot = True,cmap = colors,linewidths = 0.
    ↪4,linecolor = 'black',cbar = False,fmt = '.2f',)
plt.title('Not_Churned Customers');
```



```
fig.tight_layout(pad = 0)
```



```
[ ]: # Visualization of the objective variable in the training and test set.
df_pct_train = y_train.value_counts().to_frame().rename(index = {0: 'no', 1:
    ↪ 'yes'})
df_pct_train = df_pct_train.rename(columns = {'y': 'count'})

labels_train = df_pct_train.index.to_list()
values_train = df_pct_train.iloc[:,0]

df_pct_test = y_test.value_counts().to_frame().rename(index = {0: 'no', 1: 'yes'})
df_pct_test = df_pct_test.rename(columns = {'y': 'count'})
```

```

labels_test = df_pct_test.index.to_list()
values_test = df_pct_test.iloc[:,0]

fig, axes = plt.subplots(1, 2, figsize = (9, 4))


def autopct_fun(abs_values):
    gen = iter(abs_values)
    return lambda pct: f"{pct:.1f}%\n({next(gen)})"

axes[0].pie(x = values_train, labels = labels_train, autopct = ↵
    ↵autopct_fun(values_train),
            wedgeprops = {'linewidth':1.5, 'edgecolor':'black'},
            textprops = {'fontsize':10, 'fontweight':'bold'})
axes[0].set_title('Train', fontsize = 10, fontweight = 'bold', color = 'black')

axes[1].pie(x = values_test, labels = labels_test, autopct = ↵
    ↵autopct_fun(values_test),
            wedgeprops = {'linewidth':1.5, 'edgecolor':'black'},
            textprops = {'fontsize':10, 'fontweight':'bold'})
axes[1].set_title('Test', fontsize = 10, fontweight = 'bold', color = 'black')

fig.tight_layout()
fig.subplots_adjust(top = 0.9)
fig.show()

```



```
[ ]: import matplotlib.pyplot as plt
import seaborn as sns

# Calculate the number of rows needed based on the number of columns
num_cols = len(continuous)
num_plots_per_row = 2
num_rows = int(np.ceil(num_cols / num_plots_per_row))

# Adjust the figure size based on the number of rows and plots per row
fig, axes = plt.subplots(num_rows, num_plots_per_row, figsize=(8, 5 * num_rows))


# Flatten the axes if there's only one row
if num_rows == 1:
    axes = axes.reshape(1, -1)

for i, column in enumerate(continuous):
    row_idx = i // num_plots_per_row
    col_idx = i % num_plots_per_row

    # Create a single axis for each variable
    ax = axes[row_idx, col_idx]

    # Create boxplots for train and test data side by side
    sns.boxplot(data=[X_train_cont[column], X_test_cont[column]], ax=ax)
    ax.set_title(column)
    ax.set_xticklabels(['Train', 'Test'])

plt.tight_layout()
plt.show()
```



```
[ ]: # with new features:
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```

# Load your dataset
# df = pd.read_csv('customer_data.csv') # Replace 'customer_data.csv' with
↳ your dataset file path


# Assuming your dataset has columns like 'age', 'churn', 'city', and 'gender'

# Define the age threshold for senior citizens
senior_age_threshold = 55

# Create a new column 'Senior Citizen' based on age
df['Senior Citizen'] = df['Age'] >= senior_age_threshold

# Create a heatmap to visualize the relationship between 'Senior Citizen' and
↳ other variables
plt.figure(figsize=(50, 50))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".5f")
plt.title('Heatmap of Correlation Matrix')
plt.show()

```

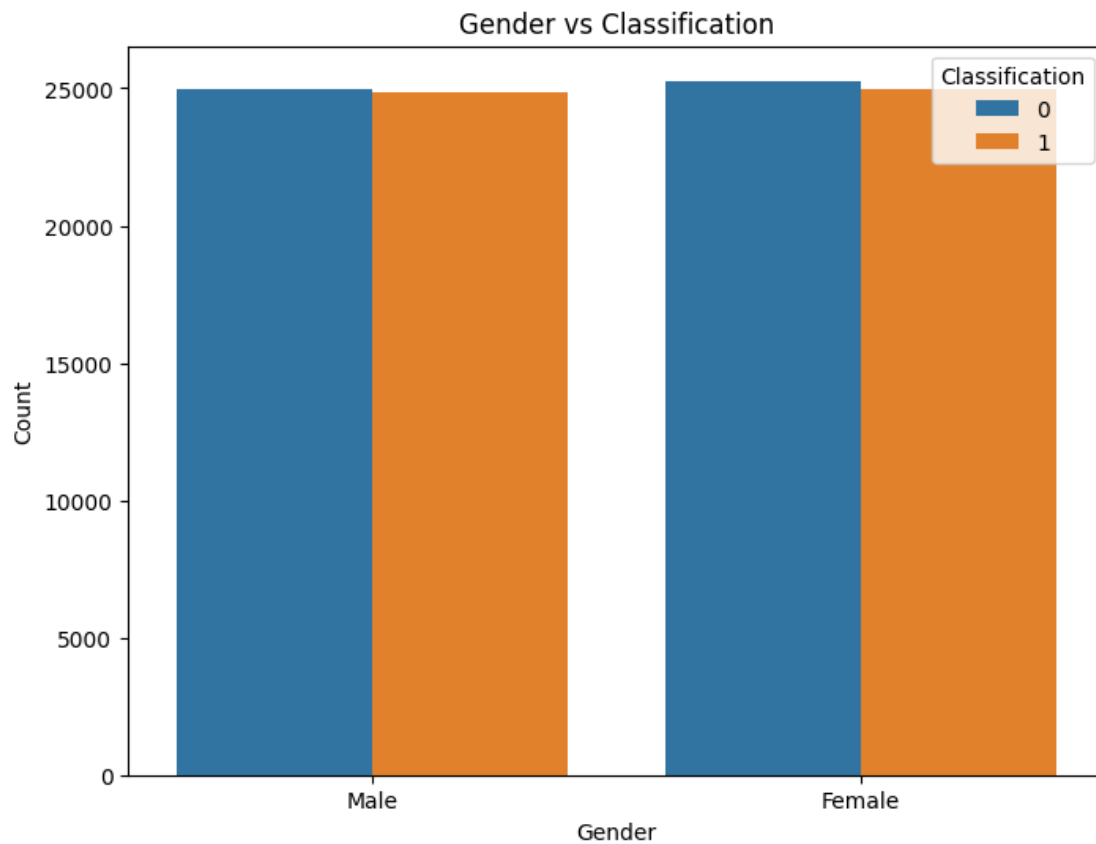


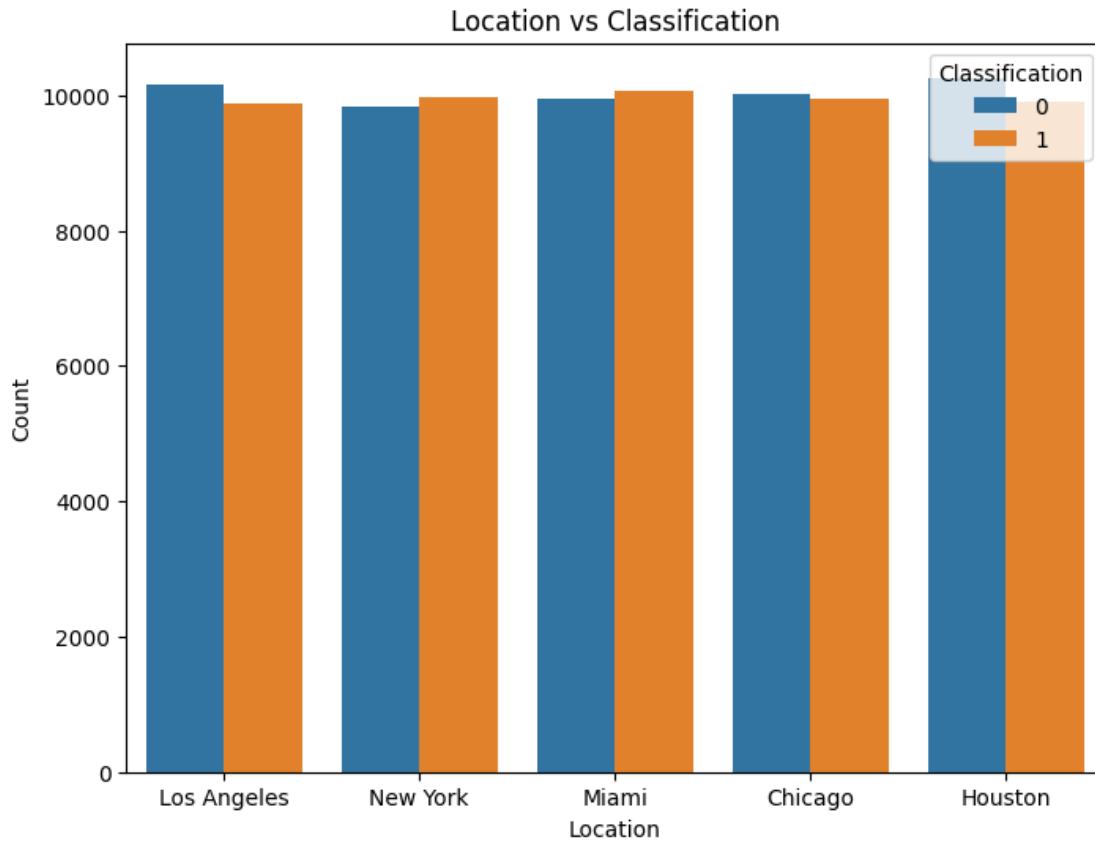
```
[ ]: import matplotlib.pyplot as plt
import seaborn as sns

# Replace 'categorical_vars' with the list of your categorical variable names
categorical_vars = ['Gender', 'Location']

# Loop through each categorical variable and create bar plots
for cat_var in categorical_vars:
```

```
plt.figure(figsize=(8, 6))
sns.countplot(x=cat_var, hue='Churn', data=df)
plt.title(f'{cat_var} vs Classification')
plt.xlabel(cat_var)
plt.ylabel('Count')
plt.legend(title='Classification', loc='upper right')
plt.show()
```





3 3 Model Building:

- Choose appropriate machine learning algorithms (e.g., logistic regression, random forest, or neural networks).
- Train and validate the selected model on the training dataset.
- Evaluate the model's performance using appropriate metrics (e.g., accuracy, precision, recall, F1-score).

```
[ ]: SEED=42
```

```
[ ]: # Estimators for the VotingClassifier and StackingClassifier.
e1 = XGBClassifier(objective = 'binary:logistic', random_state = SEED,
                  tree_method = "exact", booster = 'gbtree')
e2 = LGBMClassifier(random_state = SEED, verbose = -1)
e3 = CatBoostClassifier(random_state = SEED, verbose = 0)

voting1 = VotingClassifier(estimators = [('lgbm', e2), ('cb', e3)])
```



```
[ ]: lr = LogisticRegression(random_state = SEED)

rf = RandomForestClassifier(random_state = SEED)

et = ExtraTreesClassifier(bootstrap = True, random_state = SEED)

svc = SVC(random_state = SEED)

knn = KNeighborsClassifier()

hgb = HistGradientBoostingClassifier(random_state = SEED)

xgb = XGBClassifier(objective = 'binary:logistic', tree_method = "exact",
                    random_state = SEED, booster = 'gbtree')

lgbm = LGBMClassifier(random_state = SEED, verbose = -1)

cb = CatBoostClassifier(random_state = SEED, verbose = 0)

voting = VotingClassifier(estimators = [('lgbm', e2), ('cb', e3)])

stacking = StackingClassifier(estimators = [('lgbm', e2), ('cb', e3)],
                             final_estimator = voting1, stack_method='predict')

# We create a dictionary where we store our models with their respective names.
# base_models = {'Logistic_Regression':lr,
#                # 'Random_Forest':rf,
#                # 'Extra_Trees':et,
#                # 'SVC':svc,
#                # 'KNN':knn,
#                # 'XGB':xgb,
#                # 'LGBM':lgbm,
#                # 'CatBoost':cb,
#                # 'Voting':voting,
#                # 'Stacking':stacking,
#                # 'HistGrad':hgb}

base_models = {'Logistic_Regression':lr,
               'Random_Forest':rf,
               'Extra_Trees':et,
               'KNN':knn,
               'XGB':xgb,
               'LGBM':lgbm,
               'CatBoost':cb,
               'Voting':voting,
               'Stacking':stacking,
               'HistGrad':hgb}
```

```
[ ]: # without smote
    ## Training

    # Dictionary where we will store the metrics of each model.
    accuracy_train = {}
    accuracy_test = {}

    confusion_matrix_train = {}
    confusion_matrix_test = {}

    for model_name, model in base_models.items():
        start = time.time()
        model.fit(X_train, y_train)
        end = time.time()
        print(f'* {model_name}: {end-start} seconds')
        y_pred_train = model.predict(X_train)
        y_pred_test = model.predict(X_test)
        accuracy_train[model_name] = accuracy_score(np.array(y_train), y_pred_train)
        accuracy_test[model_name] = accuracy_score(np.array(y_test), y_pred_test)
        confusion_matrix_train[model_name] = confusion_matrix(np.array(y_train),
        ↪ y_pred_train)
        confusion_matrix_test[model_name] = confusion_matrix(np.array(y_test),
        ↪ y_pred_test)

    # We create a dataframe showing the accuracy results in training and testing.
    df_accuracy_train = pd.DataFrame.from_dict(accuracy_train, orient = 'index').
    ↪ rename(columns = {0: 'Train'})
    df_accuracy_test = pd.DataFrame.from_dict(accuracy_test, orient = 'index').
    ↪ rename(columns = {0: 'Test'})
    df_accuracy = pd.merge(df_accuracy_train, df_accuracy_test, left_index = True,
    ↪ right_index = True)
    df_accuracy = df_accuracy.sort_values(['Train', 'Test'], ascending = False)
    df_accuracy
```

```
* Logistic_Regression: 0.24571847915649414 seconds
* Random_Forest: 17.444645881652832 seconds
* Extra_Trees: 5.900312185287476 seconds
* KNN: 0.004964351654052734 seconds
* XGB: 5.007220983505249 seconds
* LGBM: 0.46331143379211426 seconds
* CatBoost: 15.28477144241333 seconds
* Voting: 13.338232040405273 seconds
* Stacking: 79.97222471237183 seconds
* HistGrad: 0.22475051879882812 seconds
```

```
[ ]:
      Train    Test
Random_Forest  1.000000  0.49240
```


Extra_Trees	1.000000	0.49180
KNN	0.686288	0.50055
XGB	0.658275	0.49855
CatBoost	0.635550	0.49765
Voting	0.619537	0.49675
Stacking	0.619537	0.49675
LGBM	0.608750	0.49635
HistGrad	0.529438	0.50355
Logistic_Regression	0.504088	0.50430

```
[ ]: # Let's visualize the metric in a bar graph.
fig,ax = plt.subplots(figsize = (20, 4.5))
n = len(df_accuracy.index)
x = np.arange(n)
width = 0.3

rects1 = ax.bar(x = x-width, height = df_accuracy.iloc[:,0], width = width)
rects2 = ax.bar(x = x, height = df_accuracy.iloc[:,1], width = width)
ax.set_xticks(x-0.12, df_accuracy.index.to_list())
ax.set_xlabel('Models', fontsize = 10, fontweight = 'bold', color = 'darkblue')
ax.set_ylabel('Score', fontsize = 10, fontweight = 'bold', color = 'darkblue')
ax.set_title('Metric of performance: Accuracy', fontsize = 12, fontweight = 'bold', color = 'darkred')

def autolabel(rects):
    for rect in rects:
        height = rect.get_height()
        ax.annotate(text = f'{round(height,4)}',
                    xy = (rect.get_x()+rect.get_width()/2, height),
                    xytext = (0,3),
                    textcoords = 'offset points',
                    ha = 'center',
                    va = 'bottom',
                    size = 6,
                    weight = 'bold',
                    color = 'black')

autolabel(rects1)
autolabel(rects2)
fig.legend(["Train", "Test"])
fig.show()
```



This time we will also try to find out the feature importance for all the features as given in the dataset after feature engineering & Normalisation

```
[ ]: X_train
```

```
[ ]: array([[0.69230769, 0.17391304, 0.77857143, ..., 0.        , 0.        ,
            1.          ],
            [0.19230769, 1.          , 0.74371429, ..., 0.        , 0.        ,
            1.          ],
            [0.75         , 0.47826087, 0.31842857, ..., 0.        , 0.        ,
            0.          ],
            ...,
            [0.98076923, 0.04347826, 0.66057143, ..., 0.        , 0.        ,
            0.          ],
            [0.71153846, 0.47826087, 0.84557143, ..., 0.        , 0.        ,
            0.          ],
            [0.15384615, 0.69565217, 0.57728571, ..., 1.        , 0.        ,
            0.          ]])
```

```
[ ]: # Performance with 3 age features
    ## Training

    # Dictionary where we will store the metrics of each model.
    accuracy_train = {}
    accuracy_test = {}

    confusion_matrix_train = {}
    confusion_matrix_test = {}

    feature_importance = {} # Dictionary to store feature importance for each model

    for model_name, model in base_models.items():
        start = time.time()
        model.fit(X_train, y_train)
        end = time.time()
        print(f'* {model_name}: {end-start} seconds')
```

```

# Calculate feature importance for the current model
if hasattr(model, 'feature_importances_'):
    feature_importance[model_name] = model.feature_importances_
else:
    feature_importance[model_name] = None

y_pred_train = model.predict(X_train)
y_pred_test = model.predict(X_test)

accuracy_train[model_name] = accuracy_score(np.array(y_train), y_pred_train)
accuracy_test[model_name] = accuracy_score(np.array(y_test), y_pred_test)

confusion_matrix_train[model_name] = confusion_matrix(np.array(y_train),
↳ y_pred_train)
confusion_matrix_test[model_name] = confusion_matrix(np.array(y_test),
↳ y_pred_test)

# We create a dataframe showing the accuracy results in training and testing.
df_accuracy_train = pd.DataFrame.from_dict(accuracy_train, orient='index').
↳ rename(columns={0: 'Train'})
df_accuracy_test = pd.DataFrame.from_dict(accuracy_test, orient='index').
↳ rename(columns={0: 'Test'})
df_accuracy = pd.merge(df_accuracy_train, df_accuracy_test, left_index=True,
↳ right_index=True)
df_accuracy = df_accuracy.sort_values(['Train', 'Test'], ascending=False)
df_accuracy

# Display feature importance for each model
for model_name, importance_scores in feature_importance.items():
    print(f"Feature Importance for {model_name}:\n{importance_scores}\n")

```

```

* Logistic_Regression: 0.24431657791137695 seconds
* Random_Forest: 14.356007099151611 seconds
* Extra_Trees: 5.80819034576416 seconds
* KNN: 0.14777493476867676 seconds
* XGB: 6.133852005004883 seconds
* LGBM: 0.37482261657714844 seconds
* CatBoost: 12.685113668441772 seconds
* Voting: 12.999969005584717 seconds
* Stacking: 80.21916794776917 seconds
* HistGrad: 0.3250269889831543 seconds
Feature Importance for Logistic_Regression:
None

```

```

Feature Importance for Random_Forest:
[0.16646724 0.16392413 0.3267288 0.31016322 0.00227566 0.00193949

```

```
0.00284723 0.00350085 0.00354699 0.00393858 0.00358155 0.00402646
0.00359787 0.00346193]
```

Feature Importance for Extra_Trees:

```
[0.20060919 0.19896479 0.29138996 0.28682991 0.00209836 0.00186667
0.00259327 0.00184835 0.00186831 0.00240161 0.00251508 0.00243365
0.00235886 0.00222201]
```

Feature Importance for KNN:

None

Feature Importance for XGB:

```
[0.09142012 0.08974636 0.09829018 0.09034033 0.          0.
0.10562229 0.08631585 0.          0.08588835 0.08964626 0.09026954
0.08443113 0.08802893]
```

Feature Importance for LGBM:

```
[636 435 775 823  0  1 17 90 15 35 39 33 51 50]
```

Feature Importance for CatBoost:

```
[17.96148895 15.61215413 24.49617662 24.86973519 0.42153705 0.59495148
1.50227705 2.23542685 1.32647558 2.08496323 2.00419713 2.64969369
1.93182017 2.30910288]
```

Feature Importance for Voting:

None

Feature Importance for Stacking:

None

Feature Importance for HistGrad:

None

```
[ ]: df_accuracy
```

```
[ ]:
      Train  Test
Extra_Trees 1.000000 0.49590
Random_Forest 1.000000 0.49355
KNN 0.687137 0.49800
XGB 0.656438 0.50260
CatBoost 0.634200 0.49660
Voting 0.619263 0.49570
Stacking 0.619263 0.49570
LGBM 0.611463 0.49935
HistGrad 0.530438 0.50395
Logistic_Regression 0.504062 0.50425
```

Features like Senior Citizen, Junior Citizen, Median Citizen, Gender_Male, Location_Chicago, Location_Los Angeles, Location_Miami, and Location_New York have consistently low importance scores across models. KNN Model:

```
[ ]: #Without normalisation:
# Performance with 3 age features
## Training

# Dictionary where we will store the metrics of each model.
accuracy_train = {}
accuracy_test = {}

confusion_matrix_train = {}
confusion_matrix_test = {}

feature_importance = {} # Dictionary to store feature importance for each model

for model_name, model in base_models.items():
    start = time.time()
    model.fit(X_train2, y_train2)
    end = time.time()
    print(f'* {model_name}: {end-start} seconds')

    # Calculate feature importance for the current model
    if hasattr(model, 'feature_importances_'):
        feature_importance[model_name] = model.feature_importances_
    else:
        feature_importance[model_name] = None

    y_pred_train2 = model.predict(X_train2)
    y_pred_test2 = model.predict(X_test2)

    accuracy_train[model_name] = accuracy_score(np.array(y_train2),
↪y_pred_train2)
    accuracy_test[model_name] = accuracy_score(np.array(y_test2), y_pred_test2)

    confusion_matrix_train[model_name] = confusion_matrix(np.array(y_train2),
↪y_pred_train2)
    confusion_matrix_test[model_name] = confusion_matrix(np.array(y_test2),
↪y_pred_test2)

# We create a dataframe showing the accuracy results in training and testing.
df_accuracy_train = pd.DataFrame.from_dict(accuracy_train, orient='index').
↪rename(columns={0: 'Train'})
df_accuracy_test = pd.DataFrame.from_dict(accuracy_test, orient='index').
↪rename(columns={0: 'Test'})
```

```

df_accuracy = pd.merge(df_accuracy_train, df_accuracy_test, left_index=True,
    ↪right_index=True)
df_accuracy = df_accuracy.sort_values(['Train', 'Test'], ascending=False)
df_accuracy

# Display feature importance for each model
for model_name, importance_scores in feature_importance.items():
    print(f"Feature Importance for {model_name}:\n{importance_scores}\n")

```

```

* Logistic_Regression: 0.0907135009765625 seconds
* Random_Forest: 10.954647064208984 seconds
* Extra_Trees: 5.647444009780884 seconds
* KNN: 0.15998244285583496 seconds
* XGB: 4.426905393600464 seconds
* LGBM: 0.3773632049560547 seconds
* CatBoost: 12.584888696670532 seconds
* Voting: 13.656790971755981 seconds
* Stacking: 77.42293548583984 seconds
* HistGrad: 0.17517542839050293 seconds
Feature Importance for Logistic_Regression:
None

```

```

Feature Importance for Random_Forest:
[0.16646724 0.16392413 0.3267288  0.31016322 0.00227566 0.00193949
 0.00284723 0.00350085 0.00354699 0.00393858 0.00358155 0.00402646
 0.00359787 0.00346193]

```

```

Feature Importance for Extra_Trees:
[0.20070098 0.19890828 0.29132962 0.28690111 0.00209314 0.00187971
 0.00261977 0.00184923 0.00187382 0.00237257 0.00250284 0.00243681
 0.00232927 0.00220285]

```

```

Feature Importance for KNN:
None

```

```

Feature Importance for XGB:
[0.09142012 0.08974636 0.09829018 0.09034033 0.          0.
 0.1056229  0.08631585 0.          0.08588835 0.08964626 0.09026954
 0.08443113 0.08802893]

```

```

Feature Importance for LGBM:
[629 439 827 770  0  0 20  87 16 41 37 40 42 52]

```

```

Feature Importance for CatBoost:
[17.96148895 15.61215413 24.49617662 24.86973519 0.42153705 0.59495148
 1.50227705 2.23542685 1.32647558 2.08496323 2.00419713 2.64969369
 1.93182017 2.30910288]

```


Feature Importance for Voting:
None

Feature Importance for Stacking:
None

Feature Importance for HistGrad:
None

```
[ ]: df_accuracy
```

```
[ ]:
```

	Train	Test
Extra_Trees	1.000000	0.49565
Random_Forest	1.000000	0.49450
KNN	0.687625	0.49650
XGB	0.656438	0.50255
CatBoost	0.634200	0.49660
Voting	0.616762	0.49890
Stacking	0.616762	0.49890
LGBM	0.607775	0.50085
HistGrad	0.530438	0.50395
Logistic_Regression	0.502575	0.49975

```
[ ]: # Let's visualize the metric in a bar graph.
fig,ax = plt.subplots(figsize = (20, 4.5))
n = len(df_accuracy.index)
x = np.arange(n)
width = 0.3

rects1 = ax.bar(x = x-width, height = df_accuracy.iloc[:,0], width = width)
rects2 = ax.bar(x = x, height = df_accuracy.iloc[:,1], width = width)
ax.set_xticks(x-0.12, df_accuracy.index.to_list())
ax.set_xlabel('Models', fontsize = 10, fontweight = 'bold', color = 'darkblue')
ax.set_ylabel('Score', fontsize = 10, fontweight = 'bold', color = 'darkblue')
ax.set_title('Metric of performance: Accuracy with 3 age features', fontsize = 12, fontweight = 'bold', color = 'darkred')

def autolabel(rects):
    for rect in rects:
        height = rect.get_height()
        ax.annotate(text = f'{round(height,4)}',
                    xy = (rect.get_x()+rect.get_width()/2, height),
                    xytext = (0,3),
                    textcoords = 'offset points',
                    ha = 'center',
```

```

        va = 'bottom',
        size = 6,
        weight = 'bold',
        color = 'black')

autolabel(rects1)
autolabel(rects2)
fig.legend(["Train", "Test"])
fig.show()

```

```

[ ]: #Without normalisation:
# Performance with 3 age features
## Training

# Dictionary where we will store the metrics of each model.
accuracy_train = {}
accuracy_test = {}

confusion_matrix_train = {}
confusion_matrix_test = {}

feature_importance = {} # Dictionary to store feature importance for each model

for model_name, model in base_models.items():
    start = time.time()
    model.fit(X_train2, y_train2)
    end = time.time()
    print(f'* {model_name}: {end-start} seconds')

    # Calculate feature importance for the current model
    if hasattr(model, 'feature_importances_'):
        feature_importance[model_name] = model.feature_importances_
    else:
        feature_importance[model_name] = None

    y_pred_train2 = model.predict(X_train2)
    y_pred_test2 = model.predict(X_test2)

    accuracy_train[model_name] = accuracy_score(np.array(y_train2),
↪y_pred_train2)
    accuracy_test[model_name] = accuracy_score(np.array(y_test2), y_pred_test2)

    confusion_matrix_train[model_name] = confusion_matrix(np.array(y_train2),
↪y_pred_train2)
    confusion_matrix_test[model_name] = confusion_matrix(np.array(y_test2),
↪y_pred_test2)

```

```

# We create a dataframe showing the accuracy results in training and testing.
df_accuracy_train = pd.DataFrame.from_dict(accuracy_train, orient='index').
    ↪rename(columns={0: 'Train'})
df_accuracy_test = pd.DataFrame.from_dict(accuracy_test, orient='index').
    ↪rename(columns={0: 'Test'})
df_accuracy = pd.merge(df_accuracy_train, df_accuracy_test, left_index=True,
    ↪right_index=True)
df_accuracy = df_accuracy.sort_values(['Train', 'Test'], ascending=False)
df_accuracy

# Display feature importance for each model
for model_name, importance_scores in feature_importance.items():
    print(f"Feature Importance for {model_name}:\n{importance_scores}\n")

```

```

* Logistic_Regression: 0.03347945213317871 seconds
* Random_Forest: 34.56139636039734 seconds
* Extra_Trees: 6.170707702636719 seconds
* KNN: 0.13585519790649414 seconds
* XGB: 7.86917519569397 seconds
* LGBM: 0.5224440097808838 seconds
* CatBoost: 14.47956919670105 seconds
* Voting: 14.985882997512817 seconds
* Stacking: 94.602698802948 seconds
* HistGrad: 0.25035524368286133 seconds
Feature Importance for Logistic_Regression:
None

```

```

Feature Importance for Random_Forest:
[0.08572883 0.04907891 0.10429026 0.10701157 0.10959937 0.12000839
 0.10394139 0.0491369 0.01652237 0.01273596 0.12152581 0.12042025]

```

```

Feature Importance for Extra_Trees:
[0.0984813 0.06733216 0.10163665 0.10323029 0.1012007 0.10700234
 0.10208763 0.06730873 0.01915719 0.01866966 0.10723561 0.10665776]

```

```

Feature Importance for KNN:
None

```

```

Feature Importance for XGB:
[0.10259998 0.08856104 0.10673914 0.10528548 0.10892413 0.10314061
 0.          0.          0.08181228 0.08819617 0.10636989 0.10837135]

```

```

Feature Importance for LGBM:
[329 156 386 418 357 433 0 0 49 18 451 403]

```

```

Feature Importance for CatBoost:
[11.5482856 3.89865863 8.91325583 11.80431674 11.40613182 14.05020688

```

```
4.97266225  2.29599524  1.86921901  1.32473985 13.94304471 13.97348342]
```

Feature Importance for Voting:

None

Feature Importance for Stacking:

None

Feature Importance for HistGrad:

None

```
[ ]: X.columns
```

```
[ ]: Index(['Age', 'Subscription_Length_Months', 'Monthly_Bill', 'Total_Usage_GB',
          'Average_Monthly_Data_Usage', 'Billing_Change_Rate',
          'Billing_As_Percentage', 'Customer_Tenure_Months', 'Churn_History',
          'Age_Group_Indicator', 'Remaining_Subscription_Length',
          'Average_Bill_Change'],
          dtype='object')
```

```
[ ]: df_accuracy
```

```
[ ]:
          Train    Test
Extra_Trees    1.000000  0.50405
Random_Forest    1.000000  0.50075
KNN              0.684450  0.50260
CatBoost         0.678113  0.50460
XGB              0.663963  0.50200
Voting           0.646250  0.50395
LGBM             0.624525  0.50200
HistGrad         0.538150  0.51150
Logistic_Regression 0.501763  0.50400
Stacking         0.480050  0.50200
```

WARNING: Runtime no longer has a reference to this dataframe, please re-run this cell and try again.

```
[ ]: #Without normalisation:
# Performance with 3 age features
## Training

# Dictionary where we will store the metrics of each model.
accuracy_train = {}
accuracy_test = {}

confusion_matrix_train = {}
confusion_matrix_test = {}
```

```

feature_importance = {} # Dictionary to store feature importance for each model

for model_name, model in base_models.items():
    start = time.time()
    model.fit(X_train3, y_train3)
    end = time.time()
    print(f'* {model_name}: {end-start} seconds')

    # Calculate feature importance for the current model
    if hasattr(model, 'feature_importances_'):
        feature_importance[model_name] = model.feature_importances_
    else:
        feature_importance[model_name] = None

    y_pred_train3 = model.predict(X_train3)
    y_pred_test3 = model.predict(X_test3)

    accuracy_train[model_name] = accuracy_score(np.array(y_train3),
    ↪ y_pred_train3)
    accuracy_test[model_name] = accuracy_score(np.array(y_test3), y_pred_test3)

    confusion_matrix_train[model_name] = confusion_matrix(np.array(y_train3),
    ↪ y_pred_train3)
    confusion_matrix_test[model_name] = confusion_matrix(np.array(y_test3),
    ↪ y_pred_test3)

    # We create a dataframe showing the accuracy results in training and testing.
    df_accuracy_train = pd.DataFrame.from_dict(accuracy_train, orient='index').
    ↪ rename(columns={0: 'Train'})
    df_accuracy_test = pd.DataFrame.from_dict(accuracy_test, orient='index').
    ↪ rename(columns={0: 'Test'})
    df_accuracy = pd.merge(df_accuracy_train, df_accuracy_test, left_index=True,
    ↪ right_index=True)
    df_accuracy = df_accuracy.sort_values(['Train', 'Test'], ascending=False)
    df_accuracy

    # Display feature importance for each model
    for model_name, importance_scores in feature_importance.items():
        print(f"Feature Importance for {model_name}:\n{importance_scores}\n")

```

```
[ ]: df_accuracy
```

```
[ ]:
```

	Train	Test
Random_Forest	1.000000	0.50830
Extra_Trees	1.000000	0.50020

CatBoost	0.685213	0.50165
KNN	0.684500	0.50255
XGB	0.663100	0.49615
Voting	0.655525	0.50445
LGBM	0.634988	0.50530
HistGrad	0.562588	0.50505
Stacking	0.510925	0.50200
Logistic_Regression	0.501763	0.50400

4. Model Optimization:

- Fine-tune the model parameters to improve its predictive performance.
- Explore techniques like cross-validation and hyperparameter tuning.

```
[ ]: from sklearn.experimental import enable_hist_gradient_boosting
from sklearn.ensemble import HistGradientBoostingClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score

# Define your HistGradientBoostingClassifier with default hyperparameters
hist_grad = HistGradientBoostingClassifier(random_state=42)

# Define the hyperparameter grid to search
param_grid_hist_grad = {
    'learning_rate': [0.01, 0.1, 0.2],
    'max_iter': [100, 200, 300],
    'max_depth': [3, 4, 5],
    'min_samples_leaf': [1, 2, 3],
    'max_bins': [100, 255, 512],
    'l2_regularization': [0.1, 0.01, 0.001]
}

# Create a GridSearchCV instance with cross-validation
grid_search_hist_grad = GridSearchCV(hist_grad, param_grid_hist_grad, cv=5,
    scoring='accuracy', n_jobs=-1)

# Fit the grid search to your training data
grid_search_hist_grad.fit(X_train3, y_train3)

# Get the best hyperparameters
best_params_hist_grad = grid_search_hist_grad.best_params_

# Initialize the HistGradientBoostingClassifier with the best hyperparameters
best_hist_grad = HistGradientBoostingClassifier(**best_params_hist_grad,
    random_state=42)

# Train the final model on the entire training dataset
best_hist_grad.fit(X_train3, y_train3)
```

```

# Make predictions on the validation set
y_val_pred_hist_grad = best_hist_grad.predict(X_test3)

# Evaluate the model's performance on the validation set
accuracy_hist_grad = accuracy_score(y_test3, y_val_pred_hist_grad)
print(f"HistGrad Validation Accuracy: {accuracy_hist_grad:.2f}")

```

HistGrad Validation Accuracy: 0.50

```

[ ]: from sklearn.neighbors import KNeighborsClassifier

# Define your K-Nearest Neighbors (KNN) Classifier with default hyperparameters
knn = KNeighborsClassifier()

# Define the hyperparameter grid to search
param_grid_knn = {
    'n_neighbors': [3, 5, 7, 9], # Adjust the number of neighbors to test
    'weights': ['uniform', 'distance'], # Weighting scheme
    'p': [1, 2] # Minkowski distance power (1 for Manhattan, 2 for Euclidean)
}

# Create a GridSearchCV instance with cross-validation
grid_search_knn = GridSearchCV(knn, param_grid_knn, cv=5, scoring='accuracy',
    ↪n_jobs=-1)

# Fit the grid search to your training data
grid_search_knn.fit(X_train3, y_train3)

# Get the best hyperparameters
best_params_knn = grid_search_knn.best_params_

# Initialize the K-Nearest Neighbors (KNN) model with the best hyperparameters
best_knn = KNeighborsClassifier(**best_params_knn)

# Train the final KNN model on the entire training dataset
best_knn.fit(X_train3, y_train3)

# Make predictions on the validation set
y_val_pred_knn = best_knn.predict(X_test3)

# Evaluate the KNN model's performance on the validation set
accuracy_knn = accuracy_score(y_test3, y_val_pred_knn)
print(f"KNN Validation Accuracy: {accuracy_knn:.2f}")

```

KNN Validation Accuracy: 0.50

```
[ ]: # lets try my xgb method here:
# X=X.drop('user-definedlabeln',axis=1)
# y=one_hot_encoded_data['user-definedlabeln']
# X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.
↳30,random_state=30)

gbm_param_grid = {
    'n_estimators': range(10, 100,10),
    'max_depth': range(3, 5),
    'learning_rate': [.4, .45, .5, .55, .6],
    'colsample_bytree': [.6, .7, .8, .9, 1]
}
gbm = XGBClassifier(n_estimators=100)
xgb_random = RandomizedSearchCV(param_distributions=gbm_param_grid,
                                estimator = gbm, scoring = "accuracy",
                                verbose = 1, n_iter = 50, cv = 5)

xgb_random.fit(X_normalized, y)
rfc3=xgb_random.predict(X_test)
# Print the best parameters and lowest RMSE
print("Best parameters found: ", xgb_random.best_params_,end="\n\n")
print("Best accuracy found: ", xgb_random.best_score_,end="\n\n")

print("accuracy=",accuracy_score(y_test,rfc3),end="\n\n")
print("confusion matrix:\n\n",classification_report(y_test,rfc3,digits=5))
```

```
[ ]: # import pandas as pd

# Assuming you have a DataFrame 'df' with the original features: Age,
↳Subscription_Length_Months, Monthly_Bill, Total_Usage_GB
# Replace 'your_dataset.csv' with the path to your dataset
# df = pd.read_csv('your_dataset.csv')

# Feature Engineering

# 1. Average Monthly Data Usage
df['Average_Monthly_Data_Usage'] = df['Total_Usage_GB'] /
↳df['Subscription_Length_Months']

# 2. Billing Change Rate
df['Billing_Change_Rate'] = df['Monthly_Bill'].diff()

# 3. Billing Amount as a Percentage
df['Billing_As_Percentage'] = (df['Monthly_Bill'] / df['Monthly_Bill'].mean())
↳
↳* 100

# 4. Customer Tenure in Months
```



```

df['Customer_Tenure_Months'] = df['Subscription_Length_Months']

# 5. Churn History (Assuming 'Churn' is a binary column indicating churn
↳ history)
df['Churn_History'] = df['Churn'].shift(1) # Lagged version of the churn column

# 6. Age Group Indicator (Assuming age groups are defined)
age_bins = [0, 30, 50, 100] # Define your age groups as needed
age_labels = ['Young', 'Middle-Aged', 'Senior']
df['Age_Group_Indicator'] = pd.cut(df['Age'], bins=age_bins, labels=age_labels)

# 7. Remaining Subscription Length
df['Remaining_Subscription_Length'] = df['Subscription_Length_Months'] - df.
↳ index

# 8. Average Bill Change
df['Average_Bill_Change'] = df['Billing_Change_Rate'].rolling(window=3).mean()

# Display the updated DataFrame with engineered features
print(df.head())

```

	CustomerID	Name	Age	Gender	Location	\
0	1	Customer_1	63	Male	Los Angeles	
1	2	Customer_2	62	Female	New York	
2	3	Customer_3	24	Female	Los Angeles	
3	4	Customer_4	36	Female	Miami	
4	5	Customer_5	46	Female	Miami	

	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	Churn	\
0	17	73.36	236	0	
1	1	48.76	172	0	
2	5	85.47	460	0	
3	3	97.94	297	1	
4	19	58.14	266	0	

	Average_Monthly_Data_Usage	Billing_Change_Rate	Billing_As_Percentage	\
0	13.882353	NaN	112.769247	
1	172.000000	-24.60	74.954041	
2	92.000000	36.71	131.384781	
3	99.000000	12.47	150.553708	
4	14.000000	-39.80	89.373010	

	Customer_Tenure_Months	Churn_History	Age_Group_Indicator	\
0	17	NaN	Senior	
1	1	0.0	Senior	
2	5	0.0	Young	
3	3	0.0	Middle-Aged	

4	19	1.0	Middle-Aged
---	----	-----	-------------

	Remaining_Subscription_Length	Average_Bill_Change
0	17	NaN
1	0	NaN
2	3	NaN
3	0	8.193333
4	15	3.126667

```
[ ]: df.tail(1)
```

```
[ ]:      CustomerID      Name  Age  Subscription_Length_Months  \
99999      100000  Customer_100000      27                      19

      Monthly_Bill  Total_Usage_GB  Churn  Average_Monthly_Data_Usage  \
99999          76.57           173      1          9.105263

      Billing_Change_Rate  Billing_As_Percentage  Customer_Tenure_Months  \
99999          27.32              117.70367                      19

      Churn_History  Age_Group_Indicator  Remaining_Subscription_Length  \
99999          1.0              2              -99980

      Average_Bill_Change  Gender_Female  Gender_Male  Location_Chicago  \
99999          4.973333              1              0              0

      Location_Houston  Location_Los Angeles  Location_Miami  \
99999              0              1              0

      Location_New York
99999              0
```

```
[ ]: import pandas as pd
from sklearn.ensemble import RandomForestClassifier

# Assuming you have engineered the additional features and added them to your
↳ dataset
# Replace 'your_dataset.csv' with the path to your dataset
# data = pd.read_csv('your_dataset.csv')

# Define your features and target variable
features = ['Age', 'Subscription_Length_Months', 'Monthly_Bill',
↳ 'Total_Usage_GB', # Original features
            'Average_Monthly_Data_Usage', 'Billing_Change_Rate',
↳ 'Billing_As_Percentage',
            'Customer_Tenure_Months', 'Churn_History', 'Age_Group_Indicator',
            'Remaining_Subscription_Length', 'Average_Bill_Change']
```

```

X = df[features] # Features
y = df['Churn'] # Target variable

# Create an imputer instance (for example, using mean imputation)

# Fit and transform the imputer on your feature data
X_imputed = imputer.fit_transform(X)

# Now, you can use X_imputed in your model

# Initialize and train a Random Forest model
model = RandomForestClassifier(random_state=42)
model.fit(X_imputed, y)

# Get feature importances
feature_importances = model.feature_importances_

# Create a DataFrame to display feature importances
importance_df = pd.DataFrame({'Feature': features, 'Importance':
    ↪feature_importances})

# Sort features by importance (highest to lowest)
importance_df = importance_df.sort_values(by='Importance', ascending=False)

# Print the feature importances
print(importance_df)

```

	Feature	Importance
10	Remaining_Subscription_Length	0.121874
5	Billing_Change_Rate	0.120788
11	Average_Bill_Change	0.120777
4	Average_Monthly_Data_Usage	0.109554
3	Total_Usage_GB	0.106813
6	Billing_As_Percentage	0.104632
2	Monthly_Bill	0.104546
0	Age	0.086257
1	Subscription_Length_Months	0.048144
7	Customer_Tenure_Months	0.047611
8	Churn_History	0.015857
9	Age_Group_Indicator	0.013146

4 4 Model Optimization:

- Fine-tune the model parameters to improve its predictive performance.
- Explore techniques like cross-validation and hyperparameter tuning.

Best Performing Model XGBoost Validation Accuracy : 0.5157

Feature Importances:

Remaining_Subscription_Length (0.1219): This feature has the highest importance score. It suggests that the remaining subscription length is a significant predictor of customer churn. Customers with shorter remaining subscription lengths might be more likely to churn.

Billing_Change_Rate (0.1208): Billing change rate is the second most important feature. It indicates that fluctuations in the monthly billing amount play a crucial role in predicting churn. Higher billing change rates might be associated with a higher likelihood of churn.

Average_Bill_Change (0.1208): This feature is very similar in importance to Billing_Change_Rate. It also captures billing fluctuations but might represent a smoothed or averaged version of those changes.

Average_Monthly_Data_Usage (0.1096): The average monthly data usage is the fourth most important feature. It suggests that customers' data usage patterns can impact their likelihood of churning. Higher data usage may correlate with lower churn rates.

Total_Usage_GB (0.1068): This feature represents the total data usage in gigabytes and is also important. It's related to Average_Monthly_Data_Usage but considers the overall data consumption.

Billing_As_Percentage (0.1046): The percentage of the monthly bill relative to some reference value is the sixth most important feature. This might capture information about customers' budget constraints or sensitivity to price changes.

Monthly_Bill (0.1045): The actual monthly billing amount is an important feature. High monthly bills could be a factor leading to churn.

Age (0.0863): Age is also significant but less so than some of the billing-related features. It indicates that customer age plays a role in predicting churn, with younger or older customers potentially having different behaviors.

Subscription_Length_Months (0.0481): Subscription length in months is less important compared to other features. However, it still contributes to the model's predictions, suggesting that longer subscription commitments might reduce churn.

Customer_Tenure_Months (0.0476): Customer tenure, or how long a customer has been with the provider, is also less important but still relevant. Longer tenure might indicate loyalty and reduce churn.

Churn_History (0.0159): Churn history, which captures whether a customer has churned in the past, has some importance. It's less critical than the other features, indicating that current behavior and characteristics are more predictive.

Age_Group_Indicator (0.0131): Age group indicator has the lowest importance in your model. It suggests that the specific age group category might not be as crucial as the overall age feature.

```
[ ]: from sklearn.experimental import enable_hist_gradient_boosting
from sklearn.ensemble import HistGradientBoostingClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score
```

```

# Define your HistGradientBoostingClassifier with default hyperparameters
hist_grad = HistGradientBoostingClassifier(random_state=42)

# Define the hyperparameter grid to search
param_grid_hist_grad = {
    'learning_rate': [0.01, 0.1, 0.2],
    'max_iter': [100, 200, 300],
    'max_depth': [3, 4, 5],
    'min_samples_leaf': [1, 2, 3],
    'max_bins': [100, 255, 512],
    'l2_regularization': [0.1, 0.01, 0.001]
}

# Create a GridSearchCV instance with cross-validation
grid_search_hist_grad = GridSearchCV(hist_grad, param_grid_hist_grad, cv=5,
    ↪scoring='accuracy', n_jobs=-1)

# Fit the grid search to your training data
grid_search_hist_grad.fit(X_train3, y_train3)

# Get the best hyperparameters
best_params_hist_grad = grid_search_hist_grad.best_params_

# Initialize the HistGradientBoostingClassifier with the best hyperparameters
best_hist_grad = HistGradientBoostingClassifier(**best_params_hist_grad,
    ↪random_state=42)

# Train the final model on the entire training dataset
best_hist_grad.fit(X_train3, y_train3)

# Make predictions on the validation set
y_val_pred_hist_grad = best_hist_grad.predict(X_test3)

# Evaluate the model's performance on the validation set
accuracy_hist_grad = accuracy_score(y_test3, y_val_pred_hist_grad)
print(f"HistGrad Validation Accuracy: {accuracy_hist_grad:.2f}")

```

```

[ ]: from sklearn.neighbors import KNeighborsClassifier

# Define your K-Nearest Neighbors (KNN) Classifier with default hyperparameters
knn = KNeighborsClassifier()

# Define the hyperparameter grid to search
param_grid_knn = {
    'n_neighbors': [3, 5, 7, 9], # Adjust the number of neighbors to test
    'weights': ['uniform', 'distance'], # Weighting scheme

```

```

    'p': [1, 2] # Minkowski distance power (1 for Manhattan, 2 for Euclidean)
}

# Create a GridSearchCV instance with cross-validation
grid_search_knn = GridSearchCV(knn, param_grid_knn, cv=5, scoring='accuracy',
    ↪n_jobs=-1)

# Fit the grid search to your training data
grid_search_knn.fit(X_train3, y_train3)

# Get the best hyperparameters
best_params_knn = grid_search_knn.best_params_

# Initialize the K-Nearest Neighbors (KNN) model with the best hyperparameters
best_knn = KNeighborsClassifier(**best_params_knn)

# Train the final KNN model on the entire training dataset
best_knn.fit(X_train3, y_train3)

# Make predictions on the validation set
y_val_pred_knn = best_knn.predict(X_test3)

# Evaluate the KNN model's performance on the validation set
accuracy_knn = accuracy_score(y_test3, y_val_pred_knn)
print(f"KNN Validation Accuracy: {accuracy_knn:.2f}")

```

```
[ ]: !pip install pycaret -q
```

```
[ ]: # pycaret trial
import pycaret.anomaly as anm
import pycaret.classification as cl
import pandas as pd
```

```
[ ]: dfsample = df.sample(frac=1)
dfsample
```

```
[ ]:
```

	Age	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	Churn	\
64111	55	22	74.13	135	1	
67350	20	18	45.55	316	1	
77239	36	18	74.98	58	1	
39551	51	19	51.63	114	0	
44014	66	24	91.28	466	0	
...	
79659	57	17	42.56	166	1	
33394	69	20	60.33	242	1	
44770	54	24	49.38	94	0	
4910	39	21	87.52	142	1	

13157	65	19	72.81	498	0
-------	----	----	-------	-----	---

	Average_Monthly_Data_Usage	Billing_Change_Rate	Billing_As_Percentage	\
64111	6.136364	-16.10	113.952893	
67350	17.555556	-54.10	70.019618	
77239	3.222222	3.88	115.259516	
39551	6.000000	16.17	79.365815	
44014	19.416667	20.81	140.315933	
...	
79659	9.764706	-10.92	65.423380	
33394	12.100000	-22.78	92.739485	
44770	3.916667	1.16	75.907107	
4910	6.761905	45.80	134.536048	
13157	26.210526	24.41	111.923785	

	Customer_Tenure_Months	Churn_History	Age_Group_Indicator	\
64111	22	0.0	1	
67350	18	0.0	2	
77239	18	1.0	0	
39551	19	1.0	1	
44014	24	0.0	1	
...	
79659	17	1.0	1	
33394	20	1.0	1	
44770	24	0.0	1	
4910	21	1.0	0	
13157	19	1.0	1	

	Remaining_Subscription_Length	Average_Bill_Change	Gender_Female	\
64111	-64089	5.356667	1	
67350	-67332	-0.523333	0	
77239	-77221	8.626667	0	
39551	-39532	4.270000	0	
44014	-43990	4.756667	0	
...	
79659	-79642	1.670000	1	
33394	-33374	5.013333	1	
44770	-44746	5.803333	1	
4910	-4889	12.240000	0	
13157	-13138	-8.413333	1	

	Gender_Male	Location_Chicago	Location_Houston	Location_Los Angeles	\
64111	0	0	0	0	
67350	1	0	0	0	
77239	1	0	0	0	
39551	1	0	0	1	
44014	1	0	1	0	

...
79659	0	1	0	0
33394	0	1	0	0
44770	0	0	0	1
4910	1	0	1	0
13157	0	0	0	1

	Location_Miami	Location_New York
64111	0	1
67350	0	1
77239	0	1
39551	0	0
44014	0	0
...
79659	0	0
33394	0	0
44770	0	0
4910	0	0
13157	0	0

[100000 rows x 20 columns]

```
[ ]: sizetest = int(0.2*(df.shape[0]))
dfTrain = dfsample[:-sizetest]
dfTest = dfsample[-sizetest:]

an = anm.setup(dfTrain)
model = anm.create_model('histogram')
preds = anm.assign_model(model)
preds = preds.drop('Anomaly',axis=1)
preds
```

<pandas.io.formats.style.Styler at 0x7e0ca6714cd0>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Processing: 0%| | 0/3 [00:00<?, ?it/s]

<IPython.core.display.HTML object>

```
[ ]:      Age  Subscription_Length_Months  Monthly_Bill  Total_Usage_GB  Churn  \
64111    55                22      74.129997         135      1
67350    20                18      45.549999         316      1
77239    36                18      74.980003          58      1
39551    51                19      51.630001         114      0
44014    66                24      91.279999         466      0
...     ...                ...                ...     ...     ...
```


93562	68	11	47.500000	452	1
49046	39	1	42.970001	52	0
18891	55	13	39.650002	170	0
20680	44	7	35.930000	150	1
71895	56	12	59.110001	172	0

	Average_Monthly_Data_Usage	Billing_Change_Rate	Billing_As_Percentage	\
64111	6.136364	-16.100000	113.952896	
67350	17.555555	-54.099998	70.019615	
77239	3.222222	3.880000	115.259514	
39551	6.000000	16.170000	79.365814	
44014	19.416666	20.809999	140.315933	
...	
93562	41.090908	-11.540000	73.017166	
49046	52.000000	-50.720001	66.053635	
18891	13.076923	-43.849998	60.950119	
20680	21.428572	-58.310001	55.231720	
71895	14.333333	-28.870001	90.864098	

	Customer_Tenure_Months	Churn_History	Age_Group_Indicator	\
64111	22	0.0	1	
67350	18	0.0	2	
77239	18	1.0	0	
39551	19	1.0	1	
44014	24	0.0	1	
...	
93562	11	1.0	1	
49046	1	0.0	0	
18891	13	1.0	1	
20680	7	0.0	0	
71895	12	1.0	1	

	Remaining_Subscription_Length	Average_Bill_Change	Gender_Female	\
64111	-64089	5.356667	1	
67350	-67332	-0.523333	0	
77239	-77221	8.626667	0	
39551	-39532	4.270000	0	
44014	-43990	4.756667	0	
...	
93562	-93551	-0.123333	1	
49046	-49045	-18.113333	1	
18891	-18878	1.946667	1	
20680	-20673	-16.766666	1	
71895	-71883	-9.160000	0	

	Gender_Male	Location_Chicago	Location_Houston	Location_Los Angeles	\
64111	0	0	0	0	

67350	1	0	0	0
77239	1	0	0	0
39551	1	0	0	1
44014	1	0	1	0
...
93562	0	1	0	0
49046	0	0	0	0
18891	0	0	1	0
20680	0	0	0	0
71895	1	0	1	0

	Location_Miami	Location_New York	Anomaly_Score
64111	0	1	11.414200
67350	0	1	8.010435
77239	0	1	7.468376
39551	0	0	11.714371
44014	0	0	11.352058
...
93562	0	0	11.615839
49046	0	1	7.544109
18891	0	0	11.751043
20680	0	1	7.774417
71895	0	0	11.810555

[80000 rows x 21 columns]

```
[ ]: classif = cl.setup(preds,target='Churn',normalize=True,
    ↪normalize_method='minmax',fold=9)
top3 = classif.compare_models(n_select=9,
    ↪fold=9,include=['ada','gbc','et','lr','rf','mlp','xgboost','lightgbm','catboost'],turbo=True)
preds
```

<pandas.io.formats.style.Styler at 0x7e0c62ce29e0>

<IPython.core.display.HTML object>

<pandas.io.formats.style.Styler at 0x7e0c62c8f130>

Processing: 0%| | 0/49 [00:00<?, ?it/s]

<IPython.core.display.HTML object>

```
[ ]:      Age  Subscription_Length_Months  Monthly_Bill  Total_Usage_GB  Churn  \
64111   55                22      74.129997         135      1
67350   20                18      45.549999         316      1
77239   36                18      74.980003          58      1
39551   51                19      51.630001         114      0
44014   66                24      91.279999         466      0
...    ...                ...                ...                ...
```

93562	68	11	47.500000	452	1
49046	39	1	42.970001	52	0
18891	55	13	39.650002	170	0
20680	44	7	35.930000	150	1
71895	56	12	59.110001	172	0

	Average_Monthly_Data_Usage	Billing_Change_Rate	Billing_As_Percentage	\
64111	6.136364	-16.100000	113.952896	
67350	17.555555	-54.099998	70.019615	
77239	3.222222	3.880000	115.259514	
39551	6.000000	16.170000	79.365814	
44014	19.416666	20.809999	140.315933	
...	
93562	41.090908	-11.540000	73.017166	
49046	52.000000	-50.720001	66.053635	
18891	13.076923	-43.849998	60.950119	
20680	21.428572	-58.310001	55.231720	
71895	14.333333	-28.870001	90.864098	

	Customer_Tenure_Months	Churn_History	...	\
64111	22	0.0	...	
67350	18	0.0	...	
77239	18	1.0	...	
39551	19	1.0	...	
44014	24	0.0	...	
...	
93562	11	1.0	...	
49046	1	0.0	...	
18891	13	1.0	...	
20680	7	0.0	...	
71895	12	1.0	...	

	Remaining_Subscription_Length	Average_Bill_Change	Gender_Female	\
64111	-64089	5.356667	1	
67350	-67332	-0.523333	0	
77239	-77221	8.626667	0	
39551	-39532	4.270000	0	
44014	-43990	4.756667	0	
...	
93562	-93551	-0.123333	1	
49046	-49045	-18.113333	1	
18891	-18878	1.946667	1	
20680	-20673	-16.766666	1	
71895	-71883	-9.160000	0	

	Gender_Male	Location_Chicago	Location_Houston	Location_Los Angeles	\
64111	0	0	0	0	

67350	1	0	0	0
77239	1	0	0	0
39551	1	0	0	1
44014	1	0	1	0
...
93562	0	1	0	0
49046	0	0	0	0
18891	0	0	1	0
20680	0	0	0	0
71895	1	0	1	0

	Location_Miami	Location_New York	Anomaly_Score
64111	0	1	11.414200
67350	0	1	8.010435
77239	0	1	7.468376
39551	0	0	11.714371
44014	0	0	11.352058
...
93562	0	0	11.615839
49046	0	1	7.544109
18891	0	0	11.751043
20680	0	1	7.774417
71895	0	0	11.810555

[80000 rows x 21 columns]

```
[ ]: an = anm.setup(dfTest)
model = anm.create_model('histogram')
dfTestsAnom = anm.assign_model(model)
dfTestsAnom = dfTestsAnom.drop('Anomaly',axis=1)
dfTestsAnom
# Description
```

<pandas.io.formats.style.Styler at 0x7e0c62d35ab0>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Processing: 0%| | 0/3 [00:00<?, ?it/s]

<IPython.core.display.HTML object>

```
[ ]: Age Subscription_Length_Months Monthly_Bill Total_Usage_GB Churn \
15177 19 2 31.990000 123 0
75601 21 8 47.070000 484 0
60617 65 11 64.629997 65 1
1425 60 17 67.180000 117 1
41658 44 12 53.570000 381 0
```

...
79659	57	17	42.560001	166
33394	69	20	60.330002	242
44770	54	24	49.380001	94
4910	39	21	87.519997	142
13157	65	19	72.809998	498

	Average_Monthly_Data_Usage	Billing_Change_Rate	Billing_As_Percentage	\
15177	61.500000	-16.510000	49.175140	
75601	60.500000	-3.480000	72.356171	
60617	5.909091	30.500000	99.349457	
1425	6.882353	22.910000	103.269333	
41658	31.750000	20.740000	82.347992	
...	
79659	9.764706	-10.920000	65.423378	
33394	12.100000	-22.780001	92.739487	
44770	3.916667	1.160000	75.907104	
4910	6.761905	45.799999	134.536041	
13157	26.210526	24.410000	111.923782	

	Customer_Tenure_Months	Churn_History	...	\
15177	2	0.0	...	
75601	8	0.0	...	
60617	11	0.0	...	
1425	17	0.0	...	
41658	12	1.0	...	
...	
79659	17	1.0	...	
33394	20	1.0	...	
44770	24	0.0	...	
4910	21	1.0	...	
13157	19	1.0	...	

	Remaining_Subscription_Length	Average_Bill_Change	Gender_Female	\
15177	-15175	-4.816667	0	
75601	-75593	-15.013333	1	
60617	-60606	-2.283333	1	
1425	-1408	8.823334	0	
41658	-41646	5.063334	0	
...	
79659	-79642	1.670000	1	
33394	-33374	5.013333	1	
44770	-44746	5.803333	1	
4910	-4889	12.240000	0	
13157	-13138	-8.413333	1	

Gender_Male	Location_Chicago	Location_Houston	Location_Los Angeles	\
-------------	------------------	------------------	----------------------	---

15177	1	0	0	0
75601	0	0	0	0
60617	0	0	0	1
1425	1	1	0	0
41658	1	0	0	1
...
79659	0	1	0	0
33394	0	1	0	0
44770	0	0	0	1
4910	1	0	1	0
13157	0	0	0	1

	Location_Miami	Location_New York	Anomaly_Score
15177	0	1	7.914317
75601	0	1	8.042049
60617	0	0	11.650193
1425	0	0	11.467817
41658	0	0	7.469678
...
79659	0	0	11.284363
33394	0	0	11.684547
44770	0	0	11.334445
4910	0	0	7.652183
13157	0	0	11.650511

[20000 rows x 21 columns]

```
[ ]: finalpreds1 = classif.predict_model(top3[0],dfTestsAnom)
finalpreds2 = classif.predict_model(top3[1],dfTestsAnom)
finalpreds3 = classif.predict_model(top3[2],dfTestsAnom)
# finalpreds
```

<pandas.io.formats.style.Styler at 0x7e0c62e1da50>

<pandas.io.formats.style.Styler at 0x7e0c62b69bd0>

<pandas.io.formats.style.Styler at 0x7e0c62c4bf70>

saving our xgboost model

```
[ ]: top3[0].save_model('/content/drive/MyDrive/internwork/xgb.bin')
```

```
[ ]: from pycaret.classification import save_model

save_model(top3[0], '/content/drive/MyDrive/internwork/xgb_caret.bin')
save_model(top3[1], '/content/drive/MyDrive/internwork/catboost_caret.bin')
save_model(top3[1], '/content/drive/MyDrive/internwork/lgbm_caret.bin')
```

Transformation Pipeline and Model Successfully Saved

Transformation Pipeline and Model Successfully Saved
Transformation Pipeline and Model Successfully Saved

```
[ ]: (Pipeline(memory=Memory(location=None),
              steps=[('numerical_imputer',
                    TransformerWrapper(exclude=None,
                                       include=['Age',
                                              'Subscription_Length_Months',
                                              'Monthly_Bill', 'Total_Usage_GB',
                                              'Average_Monthly_Data_Usage',
                                              'Billing_Change_Rate',
                                              'Billing_As_Percentage',
                                              'Customer_Tenure_Months',
                                              'Churn_History',
                                              'Age_Group_Indicator',
                                              'Remaining_Subscription_Length',
                                              'Avera...
verbose='deprecated'))),
      ('normalize',
       TransformerWrapper(exclude=None, include=None,
                          transformer=MinMaxScaler(clip=False,
                                                    copy=True,
                                                    feature_range=(0,
1))),),
      ('clean_column_names',
       TransformerWrapper(exclude=None, include=None,
                          transformer=CleanColumnNames(match='[\\]\\[\\],\\{\\}\\[\\]\\[:]+'))),
      ('trained_model',
       <catboost.core.CatBoostClassifier object at 0x7e0c62b6a1a0>)],
      verbose=False),
      '/content/drive/MyDrive/internwork/lgbm_caret.bin.pkl')
```

```
[ ]: from sklearn.metrics import confusion_matrix,accuracy_score
acc = accuracy_score(list(finalpreds['Churn'].
    ↪values),list(finalpreds['prediction_label'].astype('int').values))
cfm = confusion_matrix(list(finalpreds['Churn'].
    ↪values),list(finalpreds['prediction_label'].astype('int').values))
print("Conf. Matrix\n",cfm,"\n")
print("Accuracy: ",acc*100,"%")
```

```
[ ]: continuous=['Age', 'Subscription_Length_Months','Monthly_Bill',
    ↪'Total_Usage_GB']
categorical=[ 'Gender', 'Location']
```

```
[ ]: import matplotlib.pyplot as plt
import seaborn as sns
```

```

# Create a new DataFrame for each class of sleepiness
df_churn_0 = df[df['Churn'] == 0]
df_churn_1 = df[df['Churn'] == 1]

import matplotlib.pyplot as plt
import seaborn as sns

# Calculate the number of rows needed based on the number of columns
num_cols = len(continuous)
num_plots_per_row = 2
# num_rows = int(np.ceil(num_cols / num_plots_per_row))
num_rows = 2

# Adjust the figure size based on the number of rows and plots per row
fig, axes = plt.subplots(num_rows, num_plots_per_row, figsize=(10, 10 * num_rows))

# Flatten the axes if there's only one row
if num_rows == 1:
    axes = axes.reshape(1, -1)

for i, column in enumerate(continuous):
    row_idx = i // num_plots_per_row
    col_idx = i % num_plots_per_row

    # Create a single axis for each variable
    ax = axes[row_idx, col_idx]

    # Create boxplots for train and test data side by side
    sns.boxplot(data=[df_churn_0[column], df_churn_1[column]], ax=ax)
    ax.set_title(column)
    ax.set_xticklabels(['Churn 0', 'Churn 1'])

plt.tight_layout()
plt.show()

```

```

[ ]: X_train = pd.DataFrame(X_train, columns=X.columns)
X_train_cont = X_train[continuous] # Replace 'continous' with the actual list
    of continuous columns

X_test = pd.DataFrame(X_test, columns=X.columns)
X_test_cont = X_test[continuous] # Replace 'continous' with the actual list of
    continuous columns

```


5 5 Deployment - Done with Flask and Azure

- Once satisfied with the model's performance, deploy it into a production-like environment (you can simulate this in a development environment).
- Ensure the model can take new customer data as input and provide churn predictions.

Deployed API Link:

<https://sunbase.azurewebsites.net/predict>

POST API

Body:

Sample

```
{
  "age": 30,
  "subscription_length_months": 40,
  "monthly_bill": 76.57,
  "total_usage_gb": 173,
  "gender": "Female",
  "location": "Houston"
}
```

Python Flask Script:

This is formatted as code

Requirements.txt file:

This is formatted as code

Instructions to run locally

1. Create a folder with some name
2. Create 2 files app.py and requirements.txt
3. Run

```
pip install -r requirements.txt
```

```
python app.py
```

4. Now open a software to make a post request and use the above code as a sample